

1 Introduction

This report investigates the performance of the Simulated Annealing, and Evolutionary Strategy optimisation algorithms on the minimisation of the *Rana's function*. For each algorithm a preliminary analysis is performed on the 2D *Rana's function*, to confirm that the basis components of the algorithm are working as desired. Following this a rigorous analysis of the algorithms application to the 5D *Rana's function* for a variety of different sub-methods and hyperparameter settings is investigated, with a focus on understanding the effects of the various components of the algorithm, as well as maximising performance. For both algorithms, simple variations to the standard components were proposed, and shown to improve performance. Finally an overall comparison is performed between the 2 algorithms, using random search as a baseline. All code was written from scratch, and is available in the Appendix.

2 The Problem: *Rana's function*

The optimisation problem this report deals with is the minimisation of *Rana's function* (Whitley et al. 1996), defined by

$$\begin{aligned} & \text{Minimize} \\ & f(\mathbf{x}) = \sum_{i=1}^{n-1} x_i \cos \left(\sqrt{|x_{i+1} + x_i + 1|} \right) \sin \left(\sqrt{|x_{i+1} - x_i + 1|} \right) + \\ & \quad (1 + x_{i+1}) \cos \left(\sqrt{|x_{i+1} - x_i + 1|} \right) \sin \left(\sqrt{|x_{i+1} + x_i + 1|} \right) \\ & \text{subject to } x_i \in [-500, 500] \text{ for } i = 1, \dots, n \end{aligned} \tag{1}$$

Figure 1 shows the 2 dimensional form of *Rana's function*, showing that the problem has many local minima - making it a difficult optimisation problem. To find the best solution the following 2 elements are required: (1) a thorough exploration of the space, in order to find promising local minima, and (2) finding the lowest zone within a promising local minima. (1) and (2) have to be balanced, as they require the algorithm to perform different functions, where (1) requires testing many solutions far apart in control parameter space, and (2) requires a focuses search within a narrow zone of control parameter space. This trade-off is referred to as exploration-exploitation trade-off, and is a common theme throughout this report. Because *Rana's function* has an especially large number of local minima, it requires an especially strong emphasis on exploration.




Figure 1: *Rana's function* 2D

3 Simulated Annealing

3.1 Implementation details

For each component of the Simulated Annealing algorithm the following methods were implemented, noting any deviations from the lecture slides by inserting a "★" at the start of where they are mentioned.

Annealing schedule: Simple exponential cooling, Kirkpatrick initialisation (Kirkpatrick 1984)

Control variable perturbation methods: (1) Constant spherical noise (**Simple-step**), (2) full adaptive covariance step matrix (referred to as **Cholesky-step**, as the Cholesky decomposition is utilised), (3) adaptive diagonal covariance step matrix (**Diagonal-step**).

★ Both adaptive control variable perturbation methods are prevented from having a determinant below 10^{-16} , and each element of the step size matrix was clipped to have its absolute value within 10% of the control parameter range, to prevent the step size from becoming too small or too large in any direction, making the program more stable.

★ Significant experimentation with non-standard rules for **when** the updates to the step size control matrix occur was performed - as described in Section 3.3.2.

Convergence criteria: Absolute difference in accepted objective function $< 10^{-6}$ for last 1000 objective function evaluations. Maximum 10,000 function evaluation.

Bound enforcement: Repeated sampling of perturbation until control variable within bounds.

3.2 Model Baseline on 2D Rana function

To demonstrate that the Simulated Annealing algorithm is working as desired: (1) the basis components of the algorithm (solution generation, solution acceptance, temperature scheduling) need to be shown to be working, (2) the algorithm has to demonstrate both exploration of the space, and exploitation of the most promising local zones of the space (exploration-exploitation trade-off).

For simulated annealing to work as desired, in the initial iterations the algorithm needs to focus on exploring as much of the space as possible, while towards the end of run, instead performing a focused search of the best "zone". This is performed primarily by the temperature scheduling, which is shown in Figure 2 along with the corresponding probability of acceptances for generated solutions, over the course of a single run of the program, using the **Simple-step** method¹. As the temperature decreases, the probability of accepting solutions that increase the objective function becomes lower - focusing the optimisation on a narrower band of space. The effect of the temperature on the accepted solutions is shown in Figure 3, where the variance in the accepted objective function values is initially high - with many acceptances of objective values that were higher than the previous iteration. As the iterations continue and the temperature is annealed, both the running mean and variance of the accepted objective values generally decreases until convergence is reached. For the **Simple-step** method the step size remains fixed throughout the run, as is shown by the constant width of generated function evaluations in Figure 3. Together Figure 2 and 3 confirm that the basic mechanisms of solution generation, acceptance and temperature scheduling are working correctly together. Finally, the plot of the search pattern for this run (Figure 4) shows accepted solutions are initially distributed throughout the space, and then later on are narrowly clumped on local minima - showing that the algorithm is correctly balancing both exploration and exploitation.

¹For all initial models, a hyperparameter setting of markov-chain length = 50, annealing alpha = 0.95 is used




Figure 2: Temperature annealing and probability of acceptance for an example run of the Simulated Annealing algorithm using the `Simple-step` method.




Figure 3: Accepted and generated function evaluations for an example run of the Simulated Annealing algorithm using the `Simple-step` method.

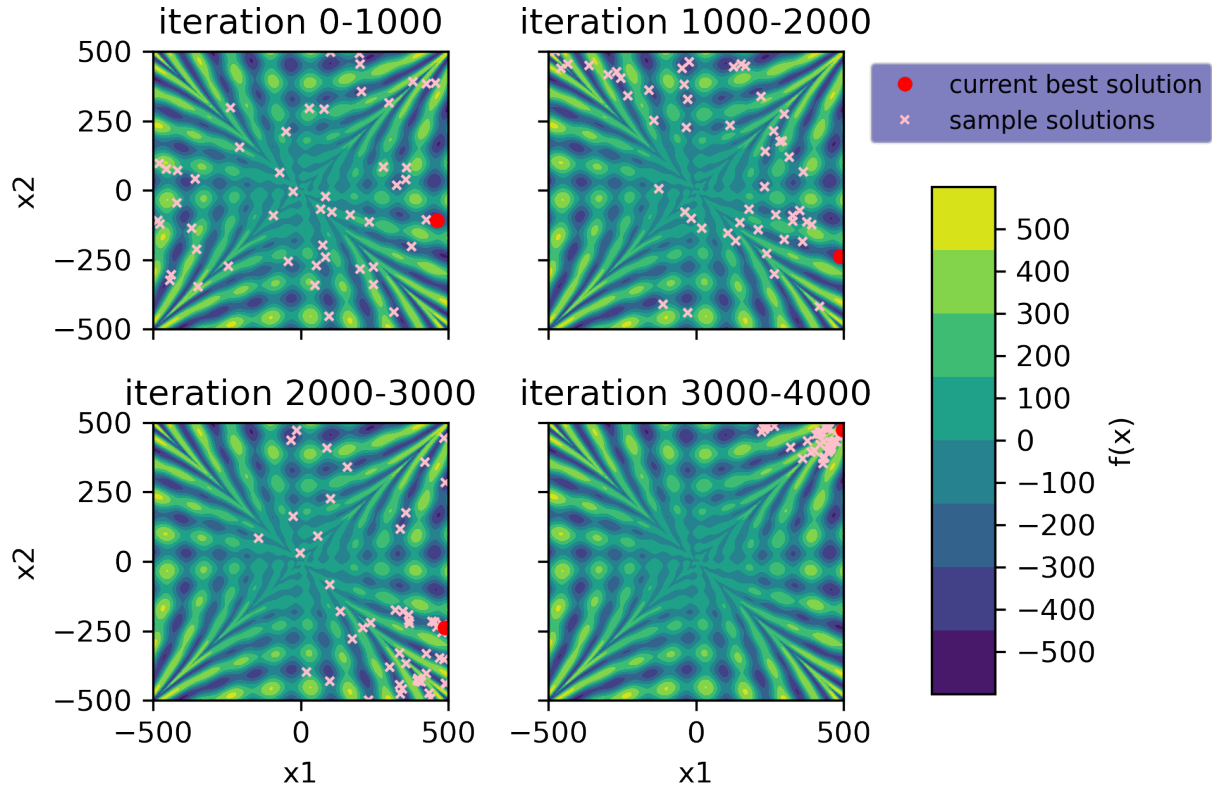


Figure 4: Sample solutions throughout optimisation for an example run of the Simulated Annealing algorithm using the **Simple-step** method. Early in the optimisation the whole space is explored. As the optimisation progresses, a progressively smaller subset of the space is focused on

The search patterns for example runs of the Simulated Annealing algorithm with the adaptive step size covariance methods (using the same seed, and temperature schedule) offer a clear visualisation of the effects of these methods, confirming that they are working as desired. Figure 5 shows the accepted solutions throughout the course of a run using the **Diagonal-step** method. Early in the run, the standard deviations of each element are high, and the temperature is high, so the accepted solutions are scattered throughout the space. The adaptability of the matrix allowing for different variance for different elements of x is then demonstrated later in the run (bottom left hand plot), where the accepted solutions have lower variance with respect to x_2 elements relative to x_1 . The example run of the **Diagonal-step** method also converges far faster than the **Simple-step** method's example run - by 2000 iterations there is virtually no variance of the accepted solutions (while with the **Simple-step** method there was still significant variance at 4000 iterations). This demonstrates ability of the **Diagonal-step** method to greatly reduce the overall step size throughout the optimisation, narrowing the space of **proposed** solutions (not just narrowing the space of accepted solutions via temperature cooling), allowing for faster convergence.

Figure 6 shows the progression of accepted solutions throughout the course of a run of the Simulated Annealing algorithm with the **Cholesky-step** method. Early in the run the additional noise from the step size covariance matrix is large and the temperature is high so the accepted solutions are scattered throughout the space. Later in the run, the accepted solutions are confined in a narrow space with a very high correlation between each of the elements of x (in the bottom left hand plot of Figure 6, the solutions appear in diagonal line). The narrowness of the space later on in the run is a function of both the lower temperature, and the lower overall noise of the sampling (from the determinant of covariance matrix decreasing). The high correlation between elements of x shows that the off-diagonal elements of the covariance in the **Cholesky-step** method are playing a significant role in the search pattern - contrasting the pattern from the **Diagonal-step** method.

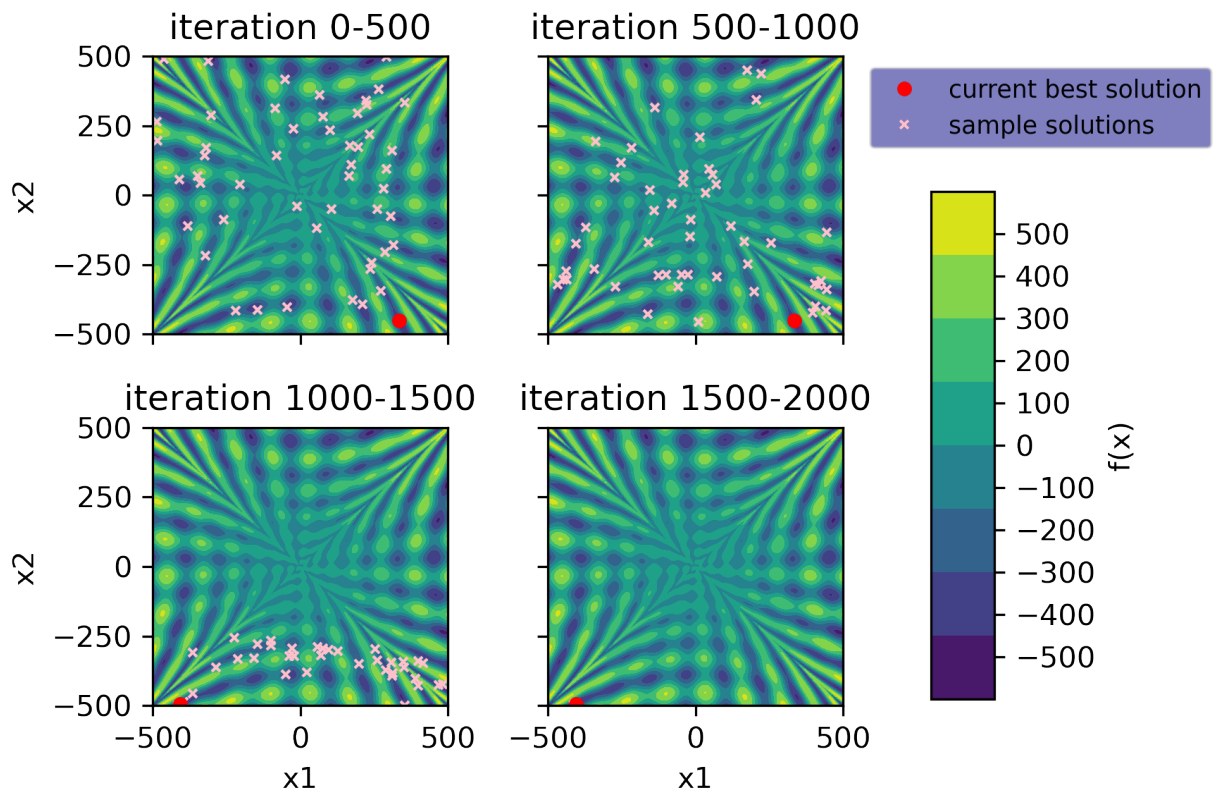


Figure 5: Sample solutions throughout optimisation for an example run of the Simulated Annealing algorithm using the **Diagonal-step** method. The difference in variance in the x_2 and x_1 elements is visible in the bottom left hand plot, where the accepted solutions have lower variance with respect to x_2 elements relative to x_1

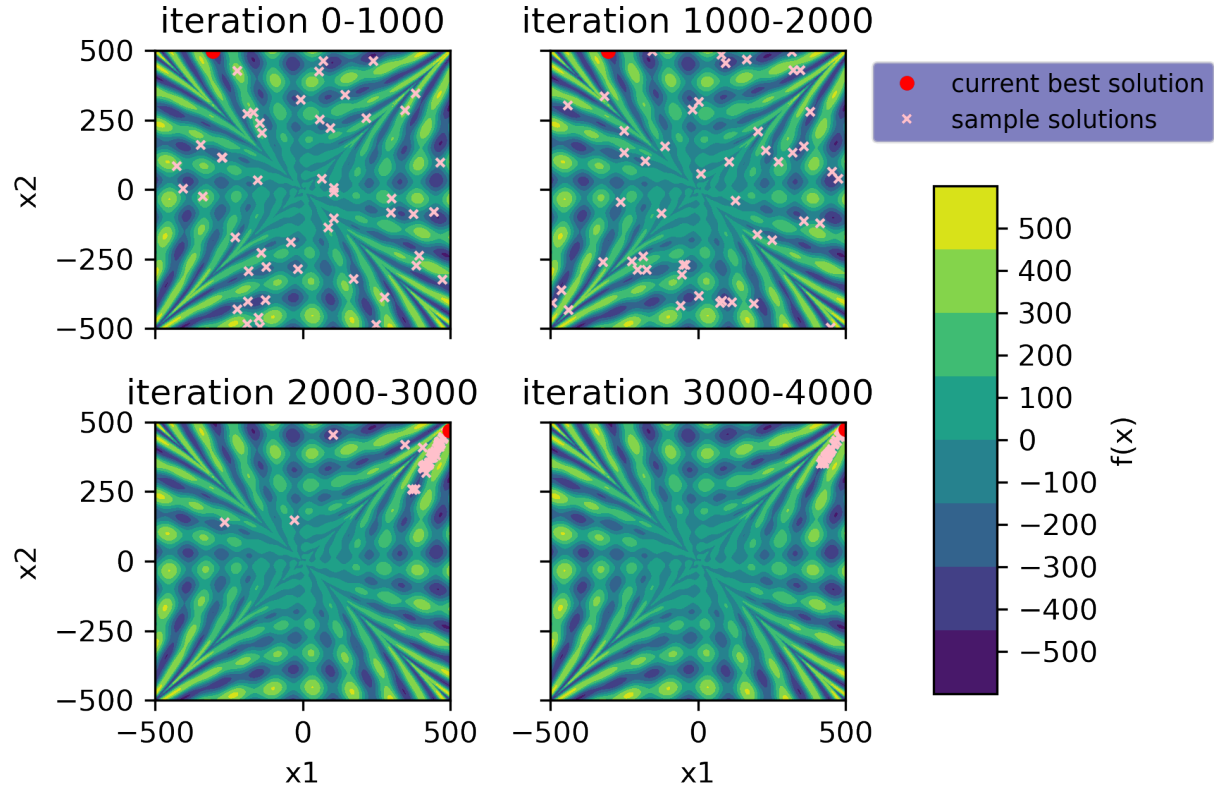


Figure 6: Sample solutions throughout optimisation for **Cholesky-step** method. The two bottom plots show significant correlation between x_1 and x_2 elements - demonstrating that the off-diagonal elements of the covariance matrix are having an effect on the search.

3.3 Analysis on 5D Rana function

In this section of the report, the effects of (1) each of the step-methods (**Simple**, **Diagonal**, **Cholesky**), (2) the step-updating rule (a proposed addition to the Simulated Annealing algorithm) and (3) the temperature annealing schedule hyperparameters (maximum markov chain length, alpha) are investigated. Each effect (1,2,3) is initially discussed and analysed separately. Because these changes to the algorithm have interacting effects, they cannot be optimised individually. Therefore a simple grid optimisation of combinations of each method and hyperparameters is performed. Finally, using the tuned hyperparameters, the overall performance of the Simulated Annealing algorithm, with the different step-size-updating rule and step-methods are compared.

3.3.1 Step Method

Various sub-plots describing the performance of the **Diagonal-step** method for an example run are given in Figure 7. The **Diagonal-step** method allows the standard deviations along each element of the control parameters to adapt over the course of the run - allowing (1) the algorithm to focus its search along some dimensions more than others and (2) the overall variance to decrease, causing the generated solutions to focus on a smaller local zone. In Figure 7, the standard deviations along each element of x fluctuate throughout the run, and correspondingly the variance in the generated function evaluations visibly varies (e.g. the standard deviation with respect to the element corresponding to the purple line increases, causing the width of the generated function evaluations to increase). By 4,000 iterations, the overall variance has decreased significantly (all of the standard deviations become very small), causing the generated objective functions to greatly decrease in variance, eventually causing the algorithm to converge the before the 10,000 maximum function evaluation count is reached. The **Diagonal-step** method has an additional advantage, which is in that it allows for the acceptance probability to be adjusted to include division by the actual step size, which scales the probability of accepting perturbations according to how large the step size was (with large step sizes having increased probability of acceptance) - thus encouraging exploration.

Various sub-plots describing the performance of the **Cholesky-step** method for a single example run, are shown in Figure 8 below. To demonstrate the effects of having a full covariance matrix (as opposed to the **Diagonal-step** method): (1) the magnitude of the eigenvalues of the covariance are plotted - indicating the magnitude of the variance along each eigenvector of the covariance matrix and (2) the minimum angle between the eigenvector corresponding to the largest eigenvalue and each of the 5 axes is plotted in order to give an indication of how much covariance between different elements of the control parameters there is.

The covariance's eigenvectors decrease throughout the program, lowering the amount of variance in x , and correspondingly in the generated objective function values (Figure 8 subplot 2 and 3). The minimum angle between the axes, and the eigenvector corresponding to the largest eigenvalue is significant throughout the course of the optimisation - this indicates that there is a significant correlation between different elements of x (i.e the non-diagonal elements of the covariance are having an effect). The **Cholesky-step** method has the advantage of potentially being able to capture a richer description of the local topology, which would allow for a more efficient "explotation" of the local zone.




Figure 7: **Diagonal-step** method example run performance. Fluctuations in the purple line cause an initial increase in the variance of the generated objective functions. The decrease in all of the standard deviations by 3000 iterations, causes the generated function evaluations to greatly decrease in variance, eventually leading in early convergence (before 10,000 iterations).

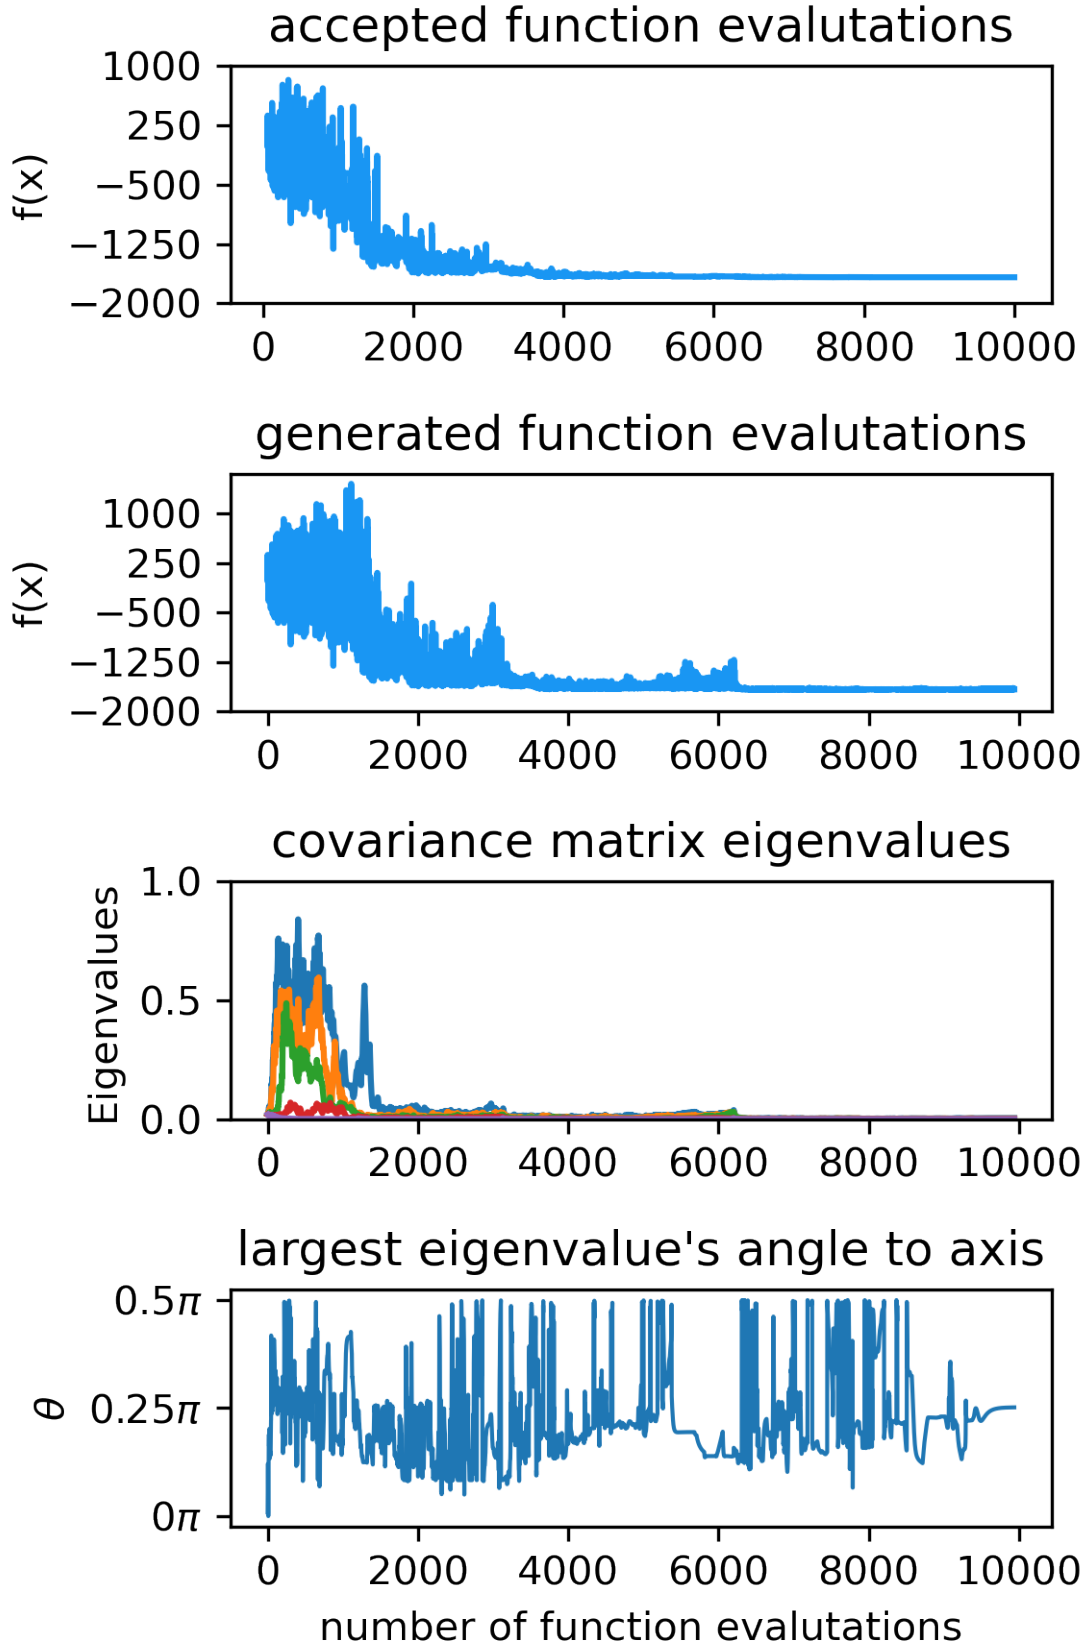


Figure 8: Cholesky-step method example performance. The magnitude of the variance in the generated objective functions has a clear link to the magnitude of the covariance eigenvectors (a decrease in the latter causing a decrease in the former). The largest eigenvalue's minimum angle to an axis is significant throughout the run, indicating that the non-diagonal elements of the covariance matrix are significant.

3.3.2 Step updating rule

For both of the adaptive step size methods the following issue was often observed: If the step size² is only updated when new values are accepted (as specified in the lecture slides, as well as in Busetti 2003), then if a reasonably good solution is found while the temperature is relatively cool and the step size is relatively large, then new solutions are almost always rejected (as the step size is too big) but because the step size is only updated when new solutions are accepted, the step size doesn't get updated, causing the program to get stuck at the current solution. I.e. there is a catch 22 where new solutions aren't accepted because the step size is too big, and the step size is not updated as there are no new solutions being accepted. An example of such a situation is shown in Figure 9, where after 2500 updates, for a good new solution to be discovered (such that it is accepted), the step size needs to be reduced, however because there are no new solutions that are good enough are discovered (because the step size is too big), the step size remains fixed and there is no improvement for the rest of the run.

This issue can be dealt with by instead updating the step size every time a perturbation is accepted **and** at an interval (e.g. every 5 steps) independent to whether the perturbation is accepted. The reason for updating the step size is to adjust the covariance matrix to better fit the local topology - and this addition to the algorithm allows for information on the local topology to be folded into the step size matrix at a rate corresponding to both the number of acceptances and the number of total steps - preventing over dependence on the rate of acceptances for the step size to match the local topology. This is referred to as the **diversified-step-update-rule** for the rest of the report.

The effect of the size of the interval between step updates using the proposed rule is shown in Figure 10. For an interval of 10,000, this method becomes identical to the original step-size update rule of only updating the step-size-matrix after acceptances. Figure 10 therefore shows that the original rule has inferior performance and that the proposed rule for updating the step size has clear benefits. A step-update-interval of 4 and 45 are preferred for the **Diagonal-step** and **Cholesky-step** methods respectively using an initial configuration for the other hyperparameter settings³⁴.




Figure 9: Illustration of issue where algorithm gets stuck, due to dependency on acceptances for step size updates.

²Here "step size" updating refers to updating the covariance matrix that controls the step size

³hyperparameter setting of markov-chain length = 50, annealing alpha = 0.95

⁴ultimately this is optimised by the grid search, so the specific results shown in Figure 10 are for illustration

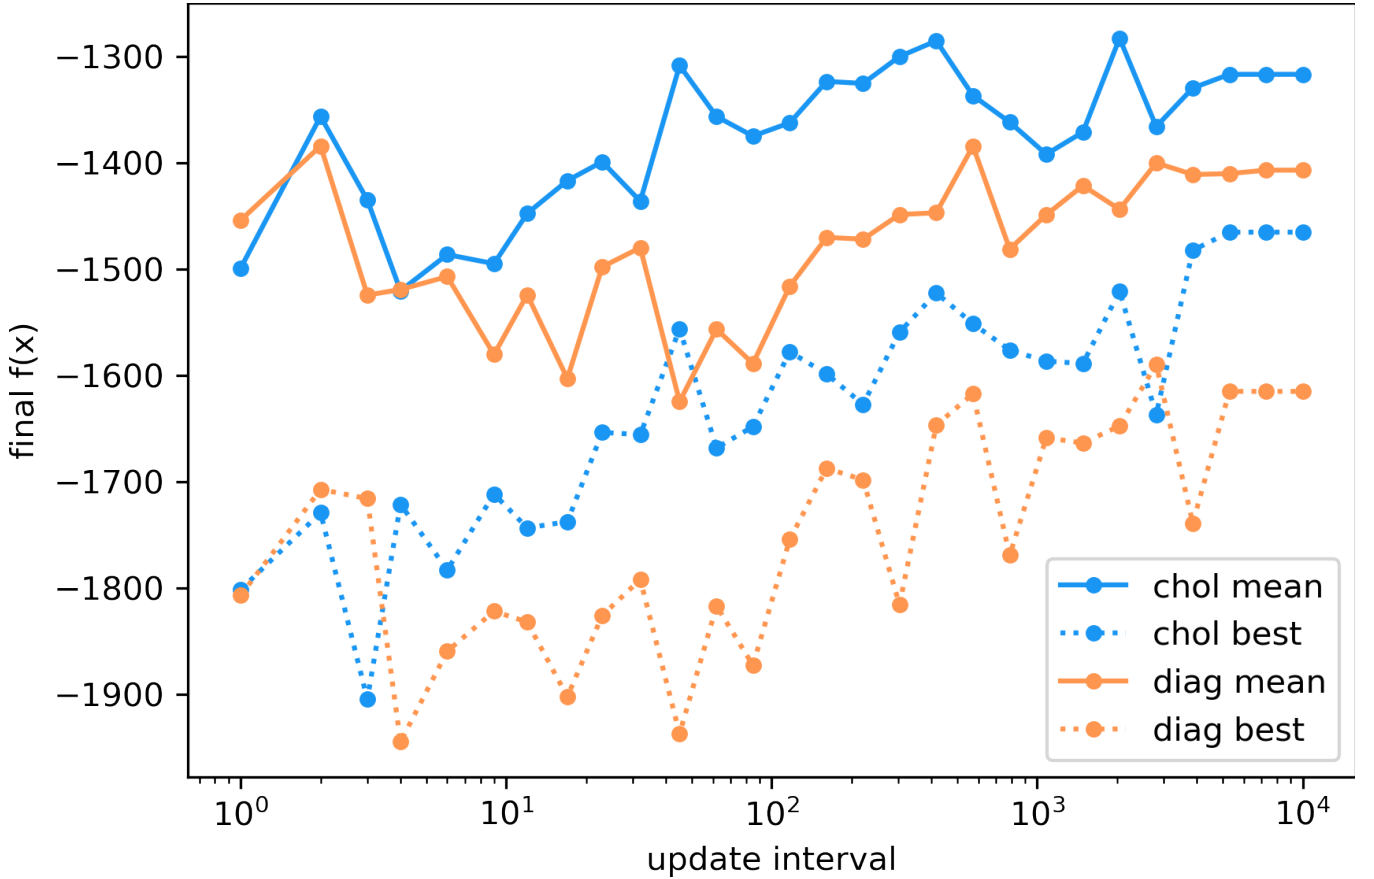


Figure 10: Effect of step-update-interval on performance. 30 runs with different random seeds are used to compute values at each point. The update interval of 10,000 is equivalent to the original step-update rule (only updating after accepted perturbations), and has clearly lower performance - showing the benefit of the **diversified-step-update-rule**

3.3.3 Temperature Schedule Hyperparameters: Markov Chain length and Alpha Parameter Optimisation

Maximum markov chain length and alpha both control the temperature schedule over the course of the Simulated Annealing algorithm. Maximum markov chain length controls the number of temperature steps, while alpha controls the size of the steps. If the markov chain length is large, then there are less total steps during the run, and thus the total temperature is reduced less. If alpha is small, the size of each temperature reduction step is large, and the temperature is reduced more over the run. This effect is shown in Figure 11. Having a well calibrated temperature schedule is important for the optimisation, as temperature controls the exploration exploitation trade-off of the program. Early in the program the temperature has to be set sufficiently high in order for the space to be explored, while later in the program the temperature has to be sufficiently low for the "best zone" to be narrowed in on. If the temperature is reduced too slowly or quickly, then the program will not sufficiently explore or exploit.

An example of the effects of maximum markov chain length and alpha on performance for **Diagonal-step** method and the **diversified-step-update-rule** are shown in Figure 12. Because both of these parameters affect the temperature schedule, there is rough equivalence in performance between certain schedules, for example (high MC length, low alpha) and (low MC length, high alpha) correspond to similar explore/exploit levels - this can be seen by the downwards sloping bands of colour in the contour plot. The average runtime is lowest for low values of both alpha and maximum markov chain length, as these lead to the system temperature becoming low quickly, resulting in convergence in lower number of steps.

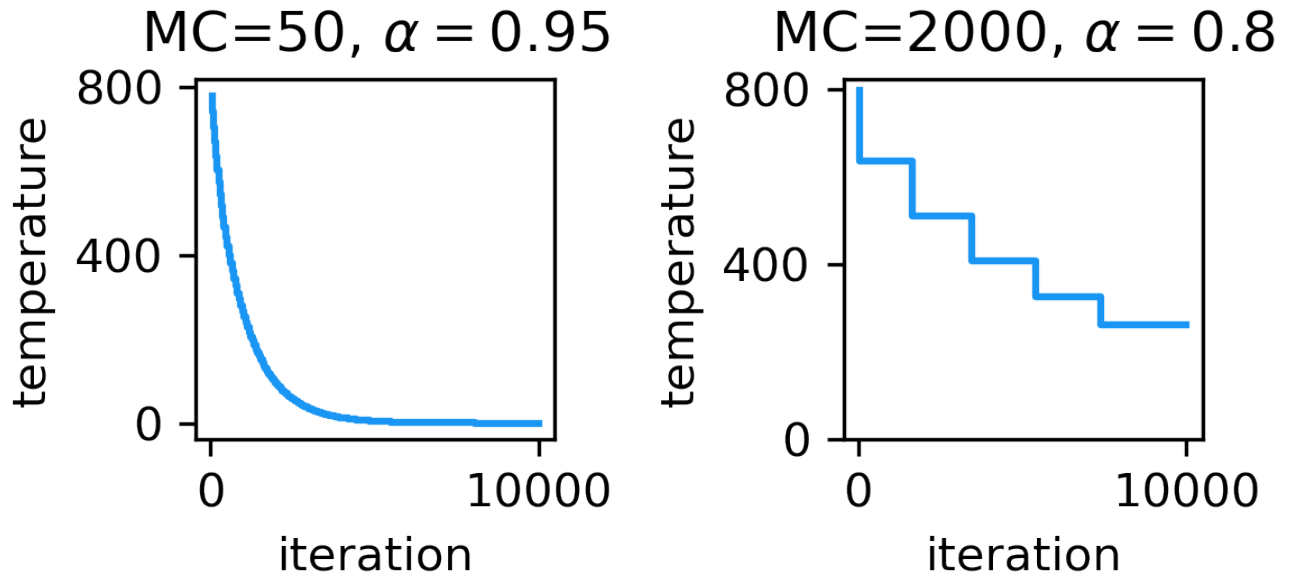


Figure 11: Effect of alpha and markov chain length (MC) on annealing schedule




Figure 12: Contour plot for grid search of maximum markov chain length, and alpha value, for **Diagonal-step** method and **diversified-step-update-rule** (interval of 20). 30 runs with 30 different random seeds are used of estimation of mean objective value and mean time

3.3.4 Results with tuned hyperparameters

A grid search was performed for each of the step-methods, varying:

1. Temperature schedule hyperparameters: Markov chain length (30 values, logarithmically spaced between 5 and 2000), alpha (30 values, linearly space between 0.5 and 0.99)
2. **diversified-step-update-rule**: Interval (30 values, linearly spaced between 1 and 60) as well as testing without the rule (i.e. using the original rule of only updating the step size after perturbation acceptances).

The results for each method, with the optimised hyperparameters, are shown in Table 1.

The **diversified-step-update-rule** has clear benefits, significantly improving both the **Cholesky-step** and **Diagonal-step** method’s mean objective function by over 1000. Both the best performing Simulated Annealing programs using the **Cholesky-step** and **Diagonal-step** methods have a lower mean objective function than the **Simple-step** method, showing the utility of the adaptive step size matrix. The **Diagonal-step** method had the best performance, both in terms of final value of the function, and in terms of runtime. Because *Rana’s function* encourages exploration, the comparative advantage of the **Cholesky-step** method to more richly match the local topology (which is useful for exploitation) did not translate into better performance compared to the **Diagonal-step** method. Furthermore, the **Diagonal-step** method included useful adjustments of the acceptance probability, incorporating the step size to encourage exploration. The **Diagonal-step** method outperforms the **Cholesky-step** in terms of runtime because it has less computationally expensive operations (i.e doesn’t have to calculate Cholesky decomposition). Although the **Diagonal-step** method requires more computation per step of the algorithm, it was also able to outperform the **Simple-step** method, by converging faster, and therefore requiring less steps.

Table 1: Results for different Simulated Annealing performance for different step-methods, with and without the diversified-step-update-rule, using tuned hyperparemeters (maximum markov chain length, alpha, and diversified-step-update-rule interval (when used)). Performance metrics calculated over 30 runs with 30 different random seeds

step-method	Cholesky		Diagonal		Simple
diversified-step-update-rule interval	-	14	-	20	-
max markov chain length	776	776	1064	9	776
annealing alpha	0.77	0.5	0.55	0.99	0.61
mean final performance	-1409.06	-1578.82	-1551.19	-1666.23	-1481.48
std dev final performance	116.75	156.9	144.45	129.67	157.47
average runtime (s)	16.13	7.98	11.4	7.23	12.13

4 Evolutionary strategies

4.1 Implementation details

For each component of the Evolutionary Strategy Algorithm, the following methods were implemented, noting any deviations from the lecture slides by inserting a "*" at the start of where they are mentioned.

Mutation methods:

1. * **simple mutation**: Spherical Gaussian noise (non-adaptive)
2. **full covariance mutation**: Gaussian noise sampled from a covariance matrix formed with standard deviation and rotation angle strategy parameters. To ensure positive-definiteness, a small diagonal matrix is added repeatedly until the covariance is positive definite. * Off-diagonal elements of the covariance matrix are clipped to have their absolute values bounded by the minimum of the diagonal elements corresponding to the row and column number (e.g. element in row 2, column 3 is clipped to have its absolute value bounded by the minimum of diagonal element 2 and diagonal element 3). This is a condition that covariance matrices hold, and helps encourage positive definiteness in a less computationally expensive manner than the original method of positive definiteness enforcement - which sometimes require repeated operations (the original method is still used in conjunction with the clipping). * Standard deviations clipped to be above 10^{-6} and below 1 to prevent mutation becoming too large or too small.
3. * **diagonal covariance mutation**: Same as **full covariance mutation** but with standard deviation strategy parameters only (i.e with rotation angles fixed at 0). Standard deviations clipped to be above 10^{-6} and below 1 to prevent mutation becoming too large or too small.

Recombination methods: Global discrete recombination for control parameters. Global intermediate recombination on strategy parameters.

Selection methods: μ, λ and $\mu + \lambda$

Convergence criteria: Absolute difference in objective function of parents $< 10^{-6}$. Maximum 10,000 iterations.

Bound enforcement: Repeated sampling of mutation until offspring within bounds.

4.2 Model Baseline on 2D Rana function

To show that the Evolution Strategy Algorithm is working as desired, an initial exploratory analysis is performed using **simple mutation** and μ, λ selection⁵, on the 2D Rana function, as it allows for visualisation of the search patterns. The key areas that are indicative of desired performance are (1) minimisation of the objective function (2) balance of explore-exploit tradeoff (3) visibility of the components of the Evolution Strategy Algorithm: Mutation, Selection and Recombination.

Figure 13 shows a sample run, in which the objective function is successfully minimised (final value -491). Mutation and crossing over are the sources of variation of objective values - and their cumulative effects can be seen by the varying mean objective values of the offspring across time. As only the best children are selected as parents, the mean objective of the parents (per generation) is significantly below the mean objective of the offspring. Figure 13 is therefore also indicative of the selection working as desired. Because **simple mutation** is used, the parents don't reach the convergence criterion (as they have non-negligible variance) even though functionally the algorithm has converged by 20 generations (the other mutation methods did in practice converge on the 2D Rana problem).

The search pattern plot (Figure 14), shows emphasis on exploration early in the run, where the solutions are spread throughout the space. The effect of selection is clear, with solutions at higher areas of the terrain becoming less prevalent as the program continues, as they are removed from the gene pool. The search pattern also shows that later in the algorithm, there is a strong focus on exploitation of the best solutions (with only 2 and 1 local zones focused on by generation 7, and 10 respectively). The symmetry of generation 4 along each axis is most likely a result of cross-over (e.g. as it allows for values of each element of various solutions to "swap", causing symmetry). Lastly, the effects of mutation are clear in the search patterns, where there are many values close (but distinctly separate) bunched at local minima.

⁵Configuration: 10 parents, 70 offspring, standard deviation of each x element = 1% of the range of x




Figure 13: Evolution strategies method example performance on 2D rana function. `simple mutation`, μ, λ selection

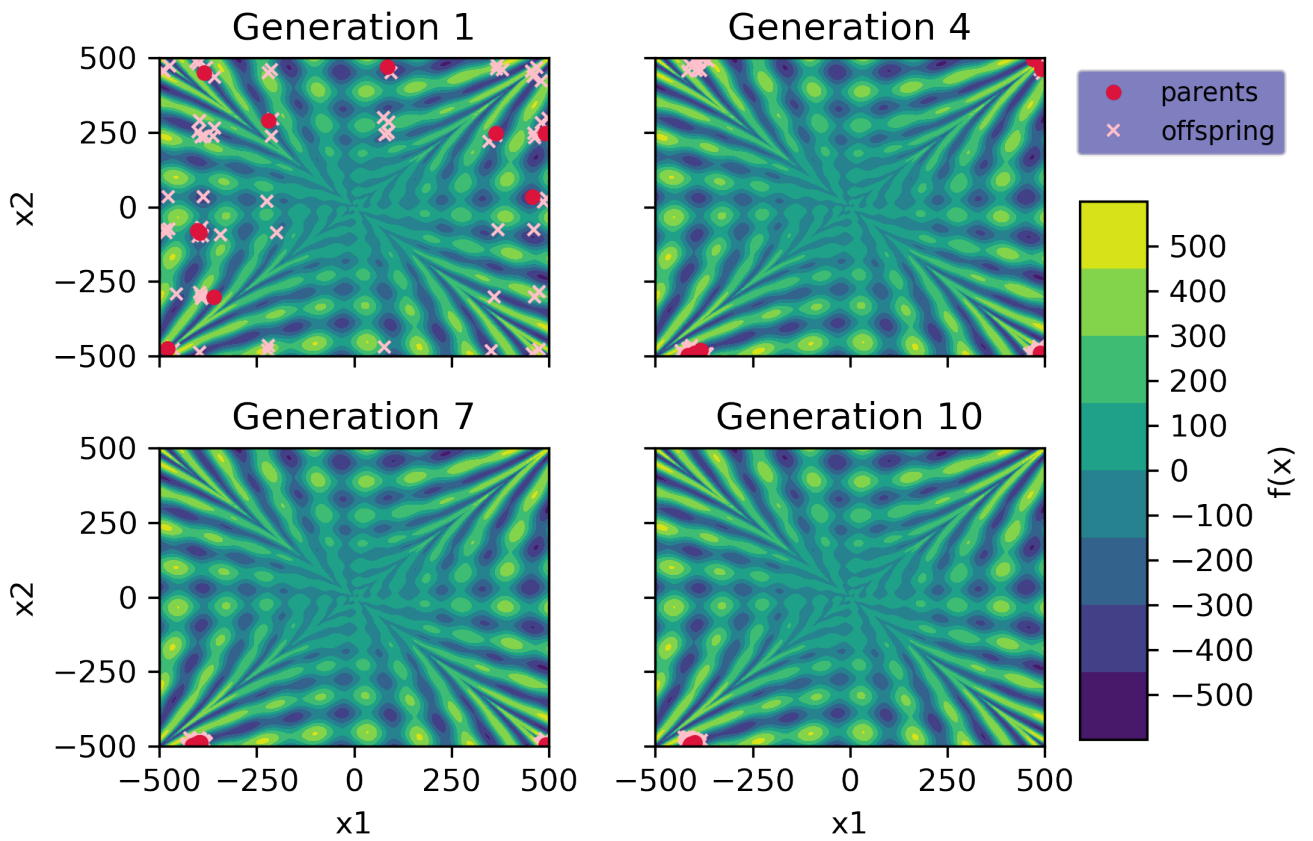


Figure 14: Evolution strategies method example search pattern on 2D rana function. `simple mutation`, μ, λ selection

4.3 Analysis on 5D Rana function

In this section, an initial presentation of the various methods for selection (μ , λ -selection and $\mu + \lambda$ -selection), and mutations (**simple mutation**, **diagonal covariance mutation**, **full covariance mutation**) is given, with an emphasis on explaining their effects (performance comparisons are rather done after hyperparameter optimisation). The effects of the hyperparameters controlling the population dynamics are explored, and optimised (by grid search) for each combination of the selection and mutation methods. Finally, a comparison of the methods with the optimal hyperparameters is given.

4.3.1 Mutation method

In evolution by natural selection, dynamic or stressful environments favour high levels of mutation, as useful adaptations are important for survival. On the other hand, stationary environments favour lower levels of mutation, as the parents typically are well suited to the environment, so children similar to the parents tend to have the highest fitness (Coyne 2010). This has a direct parallel to the Evolutionary Strategy algorithm, where the current values of the control parameters determines the "environment" (i.e. environment represented by location on control parameter landscape). Initially the control variables are in random locations that are relatively high in the landscape. As the run continues the algorithm will move significantly lower in the landscape, removing offspring similar to the original parents from the gene pool in the process. Thus the "environment" is initially dynamic, favouring high levels of mutation. Towards the end of the program, the parents are typically located in a good local minima, and change between generations becomes very small (as it is hard to find a lower point in the space). Thus later in the optimisation, the "environment" is stable, and lower levels of mutation are preferred.

The adaptive covariance methods are therefore useful as they allow the amount of mutation to adapt to the stage of the optimisation program. Furthermore, the adaptability of the magnitude mutation not only in total, but also with respect to specific elements of the control parameters (or with respect to linear combinations of control parameters in the case of **full covariance mutation**), allowing for higher levels of mutation along "directions" that are more promising.

The determinant of the covariance matrix is indicative of the total "amount" of mutation. Figure 16 shows this link, where smaller mutation-covariance determinant visibly leads to a smaller amount of variation in the offspring objective function⁶. Figure 16 shows that for both adaptive methods, the amount of mutation decreases significantly throughout the course of an example run - and thus, according to the description above, better suiting the later stationary state of the "environment".

⁶noting that selective pressure also has an effect here in reducing the variance. Comparing to Figure 13 shows that the **very** low variance comes from the lower levels of mutation (as if there is significant mutation, the standard deviation of the offspring never becomes especially small)

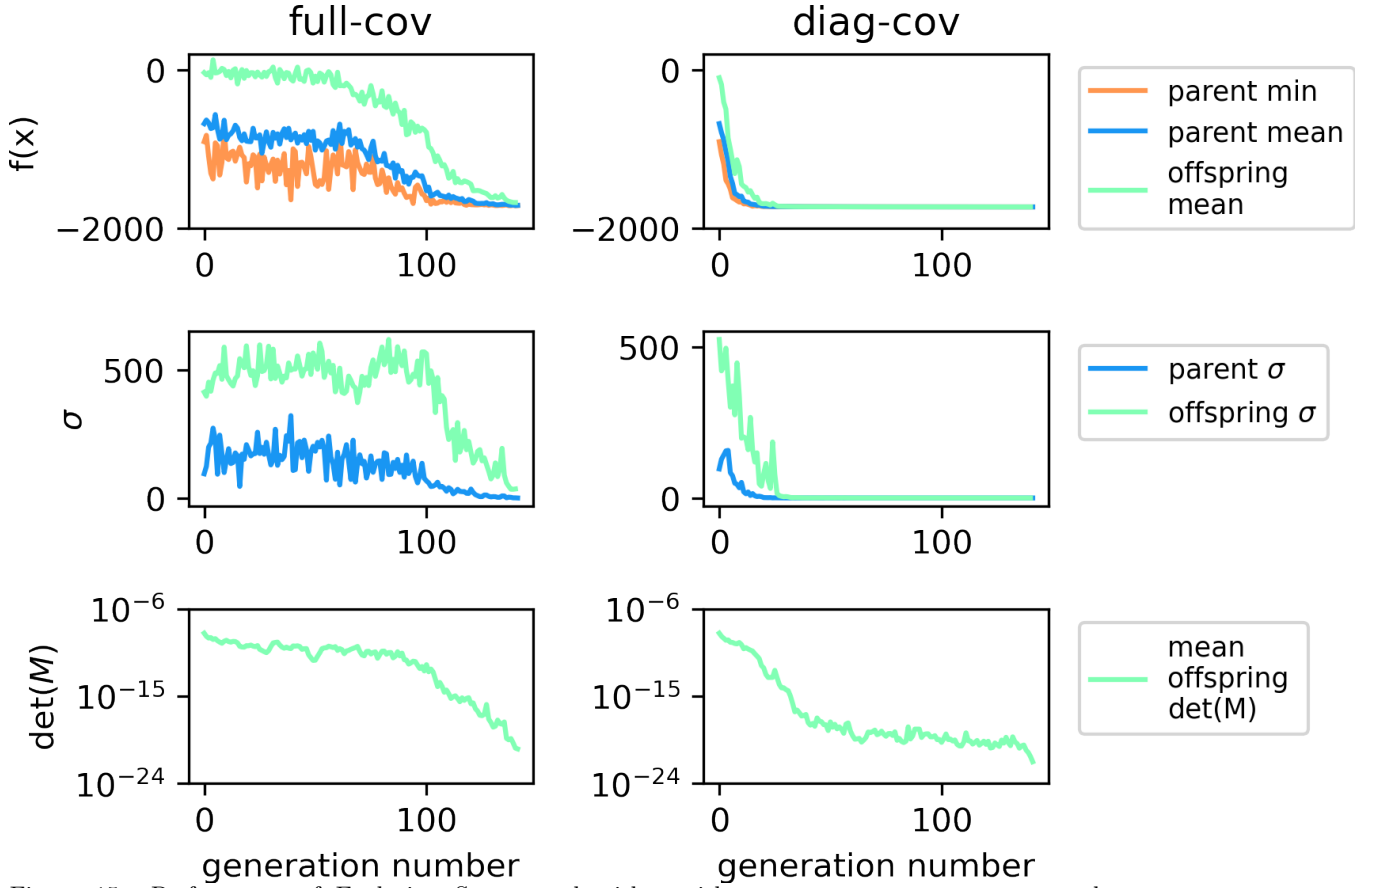


Figure 15: Performance of Evolution Strategy algorithm with full covariance mutation and diagonal covariance mutation for an example run. Note that M is the step size matrix.

4.3.2 Selection Method

μ, λ -selection is more reflective of evolution by natural selection (in which the parents of each generation die) and has the benefit of being better suited to "dynamic"⁷ environments (there is a greater amount of accumulated mutation, giving a greater probability of useful new adaption). $\mu + \lambda$ -selection has the advantage of never losing the "best" current solution, guaranteeing that the objective function will stay the same or decrease each generation - however it results in less exploration taking place, as some of the new parents in each generation are merely clones of the previous parent generation. In Figure 16, which shows an example run of both selection methods, the minimum parent of the μ, λ -selection often increases, where with $\mu + \lambda$ -selection the minimum parent objective is always lower or the same as the previous generation. In the given example run (not necessarily in general), the μ, λ -selection has a better final objective value, due to its superior exploration of the space.

⁷dynamic evolutionary environments defined in Section 4.3.1




Figure 16: Performance of Evolution Strategy algorithm with μ, λ -selection and $\mu + \lambda$ -selection

4.3.3 Population dynamics

The child to parent ratio controls the selective pressure on the system. If the child to parent ratio is high, because only a small fraction of children survive each generation, the selective pressure is high (only the very fittest survive), while if the child to parent ratio is relatively low, then many offspring survive and the selective pressure is relatively low. This relates the exploration exploitation trade-off - higher selective pressure causes a larger focus on narrow searches in the current best local optima (**exploitation**), while lower selective pressure allows for a greater **exploration** of the space. At the extreme of a 1-1 parent to child ratio, the algorithm becomes a random walk ("pure" exploration).

The effect of child to parent ratio was explored by holding the number of offspring at 100, and changing the number of parents per generation - as this allows for the number of generations (over 10,000 iterations⁸), and the amount of mutation per generation to be held constant.

The effect of child to parent ratio is demonstrated in Figure 17. The fixed-diagonal Gaussian mutation was used for this figure - as this prevents interaction effects between the mutation method and the offspring-to-parent-ratio, so allows for a simpler isolation of the effects of offspring-to-parent-ratio. In the left hand side plot, the child-to-parent-ratio is relatively low (2:1) and the following trends are clear (1) the decrease in the parents' average objective function over generations is relatively slow - from lower selective pressure creating less downwards force on the parents objectives (some parents with only mildly low objective functions are accepted each iteration). (2) The variance of the parent population is relatively high - because there are more parents, it is more likely that a parent near a different local minima to the current best will be selected. Conversely, in the right hand plot the child-to-parent-ratio is high (20:1) (1) the parents' average objective function decreases very rapidly (and then plateaus) (only the lowest objective values are selected, so the objective function quickly decreases). (2) the variance in the parents' objective functions quickly becomes very low, as the high selective pressure causes the parents to bunch around a single local optimal. Overall it is clear that the low child-to-parent ratio focuses more on exploration, and the high child to parent ratio on exploitation.

Holding the parent to child ratio and the maximum number of function evaluations constant, the population size (both parents and offspring) controls the exploration-exploitation in a different way. If the number of offspring is large then there is a lower number of total generations⁹ (as the 1000 function evaluations are "consumed" at a higher rate per generation) and if the number of offspring is small then there are many generations. As early generations focus more of exploration, and later generations on exploitation, if most of the children are contained in early generations (if there are many offspring), there is a greater focus on exploration (and correspondingly, less children \implies more exploitation). Another way this has an effect is through the number of parents - if there are many parents, then a higher number of promising local zones are passed from generation to generation. Thus for a large space with many local minima, a larger population is preferred.

⁸if convergence is reached early, then there will be less generations

⁹assuming that 10,000 function evaluations are performed (if convergence occurs early then there are less function evaluations, and less generations)

Figure 18 shows an example of the joint effects of parent to child ratio and offspring number on performance for the Evolution Strategy algorithm with **diagonal covariance** mutation and μ, λ -selection, representing a subset (**diagonal covariance** mutation and μ, λ -selection) of the grid search optimisation performed on the population dynamics hyperparameters. In practice these plots and the resultant best population dynamics hyperparameters differed significantly for different methods. This plot does show some the general trends mentioned above for the population dynamics parameters, namely (1) decreased performance if the number of children becomes too large corresponding to a lack of exploitation (2) decreased performance if the number of children is too small corresponding to a lack of exploration and (3) decreased performance if the child to parent ratio is too small, corresponding to not enough selective pressure.

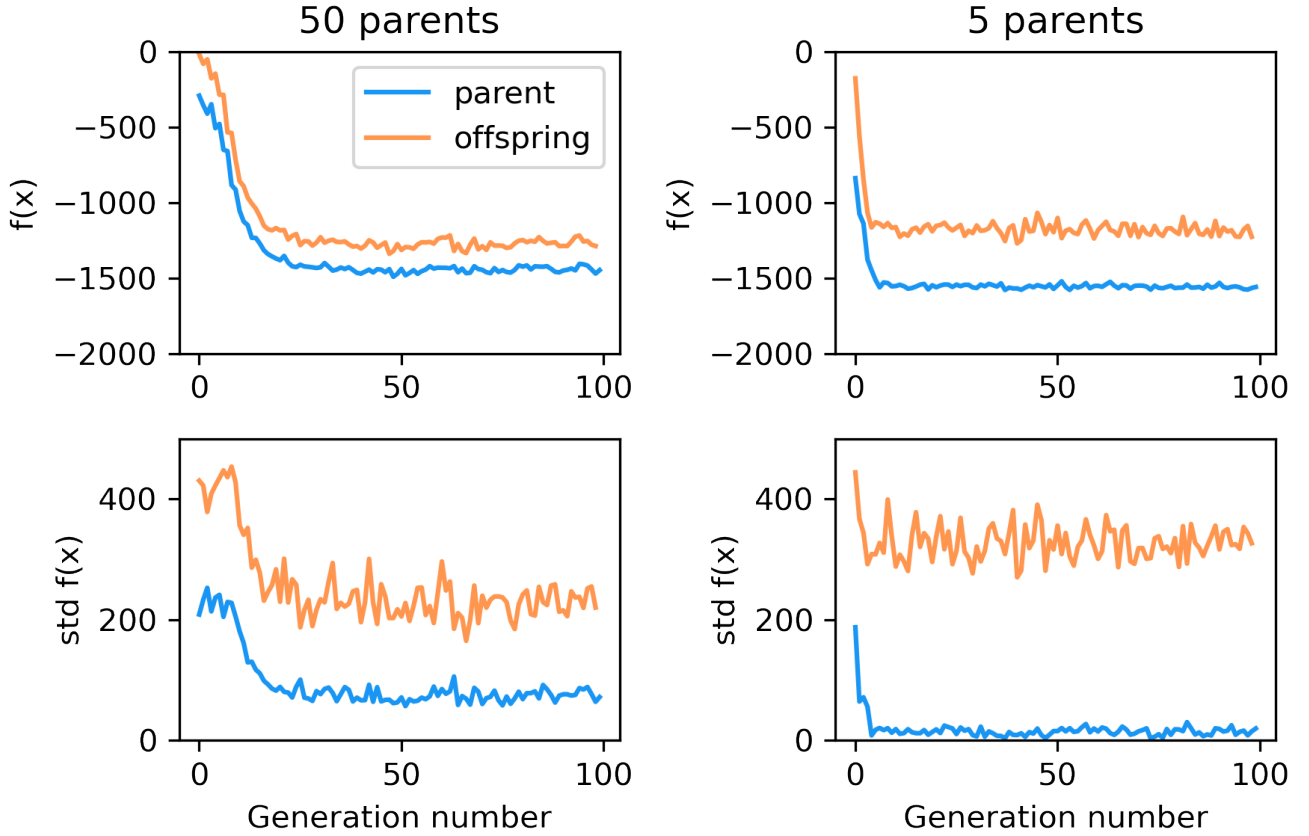


Figure 17: Illustration of the effect of child to parent ratio on performance. Number of offspring kept constant (100 offspring) Using non-adaptive mutation method, μ, λ -selection.

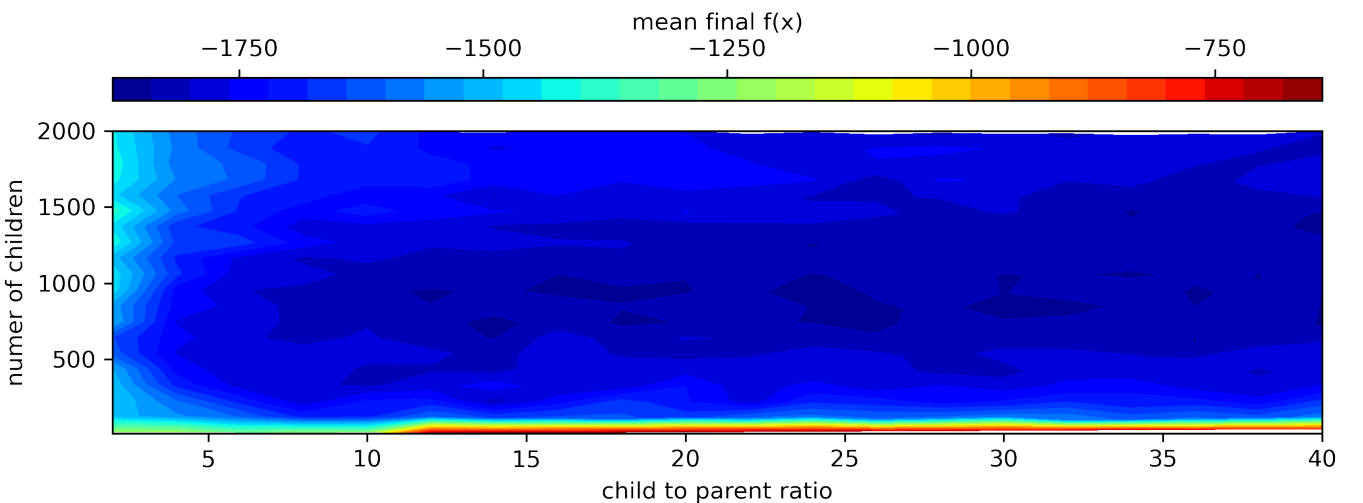


Figure 18: Contour plot of performance of the Evolution Strategy algorithm with **diagonal covariance** mutation and μ, λ -selection. Values calculated through averaging the results of 30 runs with 30 different random seeds.

4.3.4 Final Comparison with optimised population dynamics hyperparameters

A grid search optimisation was performed on the population dynamics hyperparameters for each of the mutation methods, and selection methods. The search grid was composed of 30 equally spaced values for offspring number (between 20 and 2000), and offspring per parent (between 2 and 40), with performance aggregated over 30 runs per point.

Table 2 below shows the performance results with the optimised hyperparameter settings. The **diagonal covariance mutation** method and $\mu + \lambda$ selection each generally had the best performance (with their combination yielding the best results). The **diagonal covariance mutation** method's high performance is most likely because (1) it has the benefit of an adaptive mutation, allowing it to adjust the level of mutation to fit the "evolutionary environment" (2) it has a lower number of parameters than the **full covariance mutation** and therefore is more robust. The **full covariance mutation** also has a significantly higher average runtime, due to the more expensive operation of required to (1) calculate the covariance matrix from the rotation and standard deviation strategy parameters and (2) enforce positive definiteness. $\mu + \lambda$ selection, was preferred because of the advantage of never "losing" the best solutions throughout the optimisation, and through selecting a high number of parents, the downside of $\mu + \lambda$ selection (less exploration) is minimised.

Table 2: Final results (calculated over 30 runs with 30 different random seeds), for each combination of mutation and selection methods, using optimised population dynamics hyperparameters.

mutation method	complex	complex	diagonal	diagonal	simple	simple
selection method	μ, λ	$\mu + \lambda$	μ, λ	$\mu + \lambda$	μ, λ	$\mu + \lambda$
Offspring per parent	32	14	28	40	24	40
number of offspring	96	112	952	840	936	840
parent number	3	8	34	21	39	21
mean f(x)	-1695.8	-1774	-1898.09	-1905.97	-1864.4	-1878.6
std dev f(x)	134.48	63.18	18.05	23.64	26.05	30.77
average runtime (s)	7.34	9.78	1.08	1.04	1.06	1.15

5 Overall Comparison

A comparison of the best performing Simulated Annealing and Evolution Strategy algorithms on the 5D *Rana's function* is shown in Table 3. Both algorithms are able to significantly outperform random search (which can be thought of an algorithm that focuses purely on exploration). Although this is not particularly impressive, it does provide a useful "sanity check" baseline, indicating that the algorithms are completing some functionality beyond randomly searching through the space. The Evolution Strategy algorithm is far superior to Simulated Annealing, both in terms of minimising the objective function and in terms of average runtime. The lower runtime of the Evolution Strategy Algorithm is due to it's ability to run function evaluations of each generation in parallel.

Rana's function has a roughly linear relationship between the global local minima's objective value, and the number of dimensions (Vanaret et al. 2020). Figure 19 shows the performance of the Evolution Strategy algorithm is consistently superior to random search and Simulated Annealing for an increasing number of dimensions of the *Rana's function*, with the performance difference increasing with the number of dimensions, with random search and Simulated Annealing's performance becoming quickly sub-linear. This shows that the more "difficult" the problem, the better the Evolution Strategy algorithms relative performance.

Table 3: Simulated Annealing and Evolution Strategy performance with best performing configurations, benchmarked against a random search (uniform sampling within the space for 10,000 iterations)

Method	Simulated Annealing	Evolution Strategy	Random Search
mean $f(x)$	-1666.23	-1905.97	-1498.15
std dev $f(x)$	129.67	23.64	83.08
average runtime (s)	7.23	1.04	0.27




Figure 19: Performance of Simulated Annealing, Evolution Strategy and Random Search optimisation with increasing dimension of *Rana's function*. Performances calculated for 30 runs with 30 different random seeds

References

- [1] Franco Busetti. “Simulated annealing overview”. In: (2003). URL: https://www.researchgate.net/profile/Franco_Busetti/publication/238690391_Simulated_annealing_overview/links/5c0e6c72299bf139c74de536/Simulated-annealing-overview.pdf.
- [2] Jerry A Coyne. *Why evolution is true*. Oxford University Press, 2010.
- [3] Scott Kirkpatrick. “Optimization by simulated annealing: Quantitative studies”. In: *Journal of statistical physics* 34.5-6 (1984), pp. 975–986.
- [4] Charlie Vanaret et al. “Certified global minima for a benchmark of difficult optimization problems”. In: *arXiv preprint arXiv:2003.09867* (2020).
- [5] Darrell Whitley et al. “Evaluating evolutionary algorithms”. In: *Artificial Intelligence* 85.1 (1996), pp. 245–276. ISSN: 0004-3702. DOI: [https://doi.org/10.1016/0004-3702\(95\)00124-7](https://doi.org/10.1016/0004-3702(95)00124-7). URL: <http://www.sciencedirect.com/science/article/pii/0004370295001247>.

6 Appendix

The code has been written with descriptive variable and function names, with the aim of making the steps clear without the need for excessive commenting (although comments are still given as an additional aid).

6.1 Simulated Annealing code

```
1 import numpy as np
2 from collections import deque
3
4 class SimulatedAnnealing:
5     """
6     x is transformed from the original bounds, to be bounded by (-1, 1)
7     Before outputs are given x is re-transformed back into the original form
8
9     This class includes methods (using restarts, adaptive cooling) that aren't explored
10    in the report for the sake of brevity
11    """
12    def __init__(self, x_length, x_bounds, objective_function,
13                  perturbation_method="simple",
14                  annealing_schedule="simple_exponential_cooling", with_restarts=False,
15                  archive_minimum_acceptable_dissimilarity=0.2,
16                  maximum_markov_chain_length=10, bound_enforcing_method="clipping",
17                  maximum_archive_length=20, step_size_initialisation_fraction_of_range=0.5,
18                  ↪ annealing_alpha=0.95,
19                  maximum_function_evaluations=10000,
20                  cholesky_path_length=5,
21                  perturbation_alpha = 0.1, perturbation_omega = 2.1,
22                  convergence_min_improvement = 1e-6,
23                  update_step_size_when_not_accepted_interval = 1,
24                  ):
25
26        self.x_length = x_length      # integer containing length/dimension of array x
27        self.x_bounds = x_bounds      # tuple containing the original bounds to x
28        self.x_range_internal = 2     # after being transformed to be between -1 and 1
29        self.objective_function_raw = objective_function
30        self.perturbation_method = perturbation_method
31        self.annealing_schedule = annealing_schedule
32        self.markov_chain_maximum_length = maximum_markov_chain_length
33        self.acceptances_minimum_count = round(0.6 * maximum_markov_chain_length) # 0.6 is a
34        ↪ heuristic from lectures
35        self.convergence_min_improvement = convergence_min_improvement
36        self.convergence_evaluation_window = int(maximum_function_evaluations/10)
37        # bound enforcement allows for clipping to be used (speeding up the clipping, but
38        ↪ introducing bias)
39        # the bound_enforcing_method is set to "not_clipping" throughout the report
40        self.bound_enforcing_method = bound_enforcing_method
```



```

38     # the below variable controls the "diversified-step-update-rule"
39     # i.e. how often the step-size control matrix is updated per number of iterations
40     # a value of 1 updates every step,
41     # a value of False only updates on acceptances (turns off the "step-update-rule")
42     self.update_step_size_when_not_accepted_interval =
43         ↪ update_step_size_when_not_accepted_interval
44
45     # initialise archive and parameters determining how archive is managed
46     self.archive = [] # list of (x, objective value) tuples
47     self.archive_maximum_length = maximum_archive_length
48     self.archive_minimum_acceptable_dissimilarity =
49         ↪ archive_minimum_acceptable_dissimilarity
50     self.archive_similar_dissimilarity = archive_minimum_acceptable_dissimilarity
51
52     # initialise histories and counters
53     # these are useful for inspecting the performance of the algorithm after a run
54     # and are used within the program (e.g. markov chain length)
55     self.accepted_objective_history = []
56     self.acceptances_locations = []
57     self.accepted_x_history = []
58     self.objective_history = []
59     self.x_history = []
60     self.alpha_history = []
61     self.temperature_history = []
62     self.step_size_matrix_history = []
63     self.step_size_update_locations = []
64     self.probability_of_acceptance_history = []
65     self.iterations_without_acceptance = 0
66     self.Markov_chain_length_current = 0 # initialise to 0
67     self.acceptances_count_current_chain = 0 # initialise to 0
68     self.acceptances_total_count = 0 # initialise to 0
69     self.objective_function_evaluation_count = 0 # initialise to 0
70
71     # option to restart the algorithm if no progress is made
72     self.restart_count = 0
73     self.max_iterations_without_acceptance_till_restart =
74         ↪ int(maximum_function_evaluations/10)
75     self.with_restarts = with_restarts
76
77     if perturbation_method == "simple": # step size just a constant in this case
78         self.step_size_matrix = 2*step_size_initialisation_fraction_of_range
79         # stays at initialisation in simple perturbation method
80     elif perturbation_method == "Cholesky" or "Diagonal":
81         self.perturbation_alpha = perturbation_alpha
82         self.perturbation_omega = perturbation_omega
83         self.recent_x_history = deque(maxlen=cholesky_path_length) # used for local
84         ↪ topology covariance
85     if perturbation_method == "Cholesky":
86         self.step_size_control_matrix = (np.eye(x_length) *

```

```

83         step_size_initialisation_fraction_of_range *
84         ↪ self.x_range_internal) ** 2
85     self.step_size_matrix = np.linalg.cholesky(self.step_size_control_matrix)
86 else: # Diagonal
87     self.step_size_matrix = np.eye(x_length) * \
88         step_size_initialisation_fraction_of_range *
89         ↪ self.x_range_internal
90
91 if annealing_schedule == "simple_exponential_cooling":
92     self.annealing_alpha = annealing_alpha # alpha is a constant in this case
93 elif annealing_schedule == "adaptive_cooling":
94     self.current_temperature_accepted_objective_values = [] # alpha calculated off
95     ↪ standard deviation of this
96
97 self.objective_function_evaluation_max_count = maximum_function_evaluations
98
99 def objective_function(self, x):
100     """
101     Wrapper for the objective function calls, adding some extra functionality
102     """
103     self.objective_function_evaluation_count += 1 # increment by one everytime objective
104     ↪ function is called
105     x_interp = np.interp(x, [-1, 1], self.x_bounds) # interpolation done here to pass the
106     ↪ objective function x correctly interpolated
107     result = self.objective_function_raw(x_interp)
108     return result
109
110 def run(self):
111     """
112     This function run's the major steps of the Simulated Annealing algorithm
113     # major steps in the algorithm are surrounded with a #*****
114     # other parts of this function are merely storing histories, and updating counters
115     """
116     # initialise x and temperature
117     self.x_current = self.initialise_x()
118     self.objective_current = self.objective_function(self.x_current)
119     self.initialise_temperature(self.x_current, self.objective_current)
120     self.archive.append((self.x_current, self.objective_current)) # initialise archive
121     if self.perturbation_method == "Cholesky":
122         self.recent_x_history.append(self.x_current)
123     done = False # initialise, done = True when the optimisation has completed
124     while done is False:
125         # ***** PERTUBATION *****
126         x_new = self.perturb_x(self.x_current)
127         # *****
128         objective_new = self.objective_function(x_new)

```

```

127     delta_objective = objective_new - self.objective_current
128     delta_x = x_new - self.x_current
129     self.objective_history.append(objective_new)
130     self.x_history.append(x_new)
131
132     # ***** Asses Solution *****
133     # accept change if there is an improvement, or probabilisticly (based on given
134     ↪ temperature)
135     if self.accept_x_update(delta_objective, delta_x):
136         # *****
137         self.update_archive(x_new, objective_new)
138         if self.perturbation_method == "Cholesky":
139             # store recent x values used to calculate covariance matrix
140             self.recent_x_history.append(x_new)
141         elif self.perturbation_method == "Diagonal":
142             self.latest_accepted_step = x_new - self.x_current
143
144         self.x_current = x_new # update x_current
145         self.objective_current = objective_new
146         if self.annealing_schedule == "adaptive_cooling":
147             # record accepted objective values in the current chain
148             self.current_temperature_accepted_objective_values.append(
149                 self.objective_current)
150             self.accepted_objective_history.append([objective_new,
151                 ↪ self.objective_function_evaluation_count])
152             self.acceptances_locations.append(self.objective_function_evaluation_count)
153             self.accepted_x_history.append(x_new)
154             self.acceptances_count_current_chain += 1 # in current markov chain
155             self.acceptances_total_count += 1
156             self.iterations_without_acceptance = 0
157             # *****Update Step size *****
158             self.update_step_size()
159             # *****
160         else:
161             self.iterations_without_acceptance += 1
162             # *****diversified-step-update *****
163             # update according to folding interval, using the diversified-step-update-rule
164             if self.update_step_size_when_not_accepted_interval is not False and \
165                 self.objective_function_evaluation_count %
166                 ↪ self.update_step_size_when_not_accepted_interval == 0:
167                 self.update_step_size()
168             # *****
169             self.Markov_chain_length_current += 1
170             # ***** Update Temperature *****
171             # update temperature if need be
172             # also checks for convergence
173             done = self.temperature_scheduler()
174             # *****
175             if self.with_restarts:

```

```

173         self.asses_restart()
174     if self.restart_count > 5:
175         print("max restarts reached, stopping early")
176         break
177
178     return np.interp(self.x_current, [-1, 1], self.x_bounds), self.objective_current
179
180 def accept_x_update(self, delta_objective, delta_x):
181     """
182     returns True if the latest pertubation is accepted
183     """
184     if delta_objective < 0:
185         return True
186     else:
187         if self.perturbation_method == "Diagonal":
188             probability_of_accept = np.exp(-delta_objective /
189                 ↪ (self.temperature*np.sqrt(np.sum(delta_x**2))))
189         else:
190             probability_of_accept = np.exp(-delta_objective / self.temperature)
191             self.probability_of_acceptance_history.append([np.clip(probability_of_accept, 0,
192                 ↪ 1), self.objective_function_evaluation_count])
192             if probability_of_accept > np.random.uniform(low=0, high=1):
193                 return True
194             else:
195                 return False
196
197 def initialise_x(self):
198     # initialise x randomly within the given bounds
199     return np.random.uniform(low=-1, high=1, size=self.x_length)
200
201 def initialise_temperature(self, x_current, objective_current, n_steps=60,
202     ↪ average_accept_probability=0.8):
203     """
204     Initialises system temperature using Kirkpatrick method
205     As all x's are initially accepted, x does a random walk, so changes in x are discarded
206     """
207     objective_increase_history = [] # if many samples are taken then this could be changed
208     ↪ to running average
209     for step in range(1, n_steps+1):
210         x_new = self.perturb_x(x_current)
211         objective_new = self.objective_function(x_new)
212         if objective_new > objective_current:
213             objective_increase_history.append(objective_new - objective_current)
214         if step == n_steps:
215             self.latest_accepted_step = x_new - self.x_current
216             x_current = x_new
217             objective_current = objective_new
218
219     initial_temperature = - np.mean(objective_increase_history) /
220     ↪ np.log(average_accept_probability)

```

```

218         self.temperature = initial_temperature
219         self.temperature_history.append([self.temperature, self.acceptances_total_count,
    ↪ self.objective_function_evaluation_count])
220
221     def is_positive_definite(self, matrix):
222         try:
223             np.linalg.cholesky(matrix)
224             return True
225         except:
226             return False
227
228     def update_step_size(self):
229         """
230         Update the matrix controlling the step size
231         """
232         # record when step size updates happened
233         self.step_size_update_locations.append(self.objective_function_evaluation_count)
234         if self.perturbation_method == "simple":
235             return # no update with the simple method
236         elif self.perturbation_method == "Cholesky":
237             covariance = np.cov(self.recent_x_history, rowvar=False)
238             # prevent covariance from becoming too large by clipping
239             covariance = np.clip(covariance, -self.x_range_internal / 2, self.x_range_internal
    ↪ / 2)
240             self.step_size_control_matrix = (1 - self.perturbation_alpha) *
    ↪ self.step_size_control_matrix + \
241                 self.perturbation_alpha * self.perturbation_omega *
    ↪ covariance
242             # conservative clipping preventing step size control matrix from exploding or
    ↪ getting too small
243             self.step_size_control_matrix =
    ↪ self.step_size_control_matrix.clip(-self.x_range_internal * 2,
    ↪ self.x_range_internal * 2)
244             i = 0
245             while not self.is_positive_definite(self.step_size_control_matrix):
246                 i += 1
247                 self.step_size_control_matrix += np.eye(self.x_length)*1e-6*10**i # to make
    ↪ positive definite
248                 if i > 7:
249                     raise Exception("couldn't get positive definite step size control matrix")
250             if np.linalg.det(self.step_size_matrix) < 1e-16:
251                 # increase step size if determinant falls below 1e-6
252                 self.step_size_control_matrix = self.step_size_control_matrix * 1.1
253                 # now have to enforce positive definiteness again
254                 while not self.is_positive_definite(self.step_size_control_matrix):
255                     i += 1
256                     self.step_size_control_matrix += np.eye(self.x_length) * 1e-6 * 10 ** i
257                     if i > 7:
258                         raise Exception("couldn't get positive definite step size control
    ↪ matrix")

```

```

259         self.step_size_matrix = np.linalg.cholesky(self.step_size_control_matrix)
260         self.step_size_matrix_history.append(self.step_size_matrix)
261
262     elif self.perturbation_method == "Diagonal":
263
264         self.step_size_matrix = \
265             (1-self.perturbation_alpha)*self.step_size_matrix + \
266
267             ↪ np.diag(self.perturbation_alpha*self.perturbation_omega*np.abs(self.latest_accepted_step_size))
268         # conservative clipping to prevent step size becoming too small or large
269         self.step_size_matrix = np.clip(self.step_size_matrix, self.x_range_internal *
270             ↪ 1e-16, self.x_range_internal * 2)
271
272         if np.linalg.det(self.step_size_matrix) < 1e-16:
273             # increase step size if determinant falls below 1e-6
274             self.step_size_matrix = self.step_size_matrix*1.1
275             self.step_size_matrix_history.append(np.diag(self.step_size_matrix))
276
277     def perturb_x(self, x):
278         if self.perturbation_method == "simple":
279             u_random_sample = np.random.uniform(low=-1, high=1, size=self.x_length)
280             x_new = x + self.step_size_matrix * u_random_sample # constant step size
281             if self.bound_enforcing_method == "clipping":
282                 return np.clip(x_new, -1, 1)
283             else:
284                 while max(x_new) > 1 or min(x_new) < -1:
285                     indxs_breaking_bounds = np.where((x_new > 1) + (x_new < -1) == 1)
286                     u_random_sample = np.random.uniform(low=-1, high=1,
287                         ↪ size=indxs_breaking_bounds[0].size)
288                     x_new[indxs_breaking_bounds] = x[indxs_breaking_bounds] +
289                         ↪ self.step_size_matrix * u_random_sample
290
291         elif self.perturbation_method == "Cholesky":
292             u_random_sample = np.random.uniform(low=-np.sqrt(3), high=np.sqrt(3),
293                 ↪ size=self.x_length)
294             x_new = x + self.step_size_matrix@u_random_sample
295             if self.bound_enforcing_method == "clipping":
296                 return np.clip(x_new, -1, 1)
297             else:
298                 if max(x_new) > 1 or min(x_new) < -1:
299                     x_new = self.perturb_x(x) # recursively call perturb until sampled
300                     ↪ within bounds
301
302         elif self.perturbation_method == "Diagonal":
303             u_random_sample = np.random.uniform(low=-1, high=1, size=self.x_length)
304             x_new = x+self.step_size_matrix@u_random_sample
305             if self.bound_enforcing_method == "clipping":
306                 return np.clip(x_new, -1, 1)

```

```

302         else:
303             while max(x_new) > 1 or min(x_new) < -1: # only sample specific indices not
304                 ↳ within bounds
305                 indxs_breaking_bounds = np.where((x_new > 1) + (x_new < -1) == 1)
306                 u_random_sample = np.random.uniform(low=-1, high=1,
307                 ↳ size=indxs_breaking_bounds[0].size)
308                 x_new[indxs_breaking_bounds] = \
309                     x[indxs_breaking_bounds] + \
310                     np.diag(np.diag(self.step_size_matrix)[indxs_breaking_bounds])\
311                     @u_random_sample
312
313         return x_new
314
315     def asses_restart(self, min_difference = 0.01):
316         length = self.markov_chain_maximum_length # base on length of markov chain
317         if len(self.objective_history) % length == 0 and \
318             len(self.accepted_objective_history) > self.markov_chain_maximum_length:
319             if max(self.accepted_objective_history_array[-length:, 0]) -
320                 ↳ min(self.accepted_objective_history_array[-length:, 0]) < min_difference:
321                 # then rebase from best archive solution
322                 x_restart = self.archive_x[np.argmax(self.archive_f), :]
323                 print("restarted due to minimal progress")
324                 self.restart_count += 1
325                 self.x_current = x_restart
326             elif self.iterations_without_acceptance >
327                 ↳ self.max_iterations_without_acceptance_till_restart:
328                 x_restart = self.archive_x[np.argmax(self.archive_f), :]
329                 print(f"restarted due to {self.max_iterations_without_acceptance_till_restart}
330                 ↳ iterations "
331                     f"without acceptance")
332                 self.restart_count += 1
333                 self.x_current = x_restart
334
335     def temperature_scheduler(self):
336         if self.Markov_chain_length_current > self.markov_chain_maximum_length or \
337             self.acceptances_count_current_chain > self.acceptances_minimum_count:
338             if self.annealing_schedule == "simple_exponential_cooling":
339                 self.temperature = self.temperature * self.annealing_alpha
340             elif self.annealing_schedule == "adaptive_cooling":
341                 if len(self.current_temperature_accepted_objective_values) <= 1:
342                     # if no values have been accepted, then don't change alpha
343                     # algorithm most likely close to convergence
344                     self.alpha = 1
345                     self.alpha_history.append([self.alpha,
346                     ↳ self.objective_function_evaluation_count])
347             else:
348                 latest_temperature_standard_dev =
349                 ↳ np.std(self.current_temperature_accepted_objective_values)

```



```

344         # use adaptive temperature rule
345         self.alpha = np.max([0.5,
346                               ↪ np.exp(-0.7*self.temperature/latest_temperature_standard_dev)])
347         self.alpha_history.append([self.alpha,
348                                     ↪ self.objective_function_evaluation_count])
349         self.temperature = self.temperature * self.alpha
350         if np.isnan(self.temperature):
351             self.temperature = 1e-16
352             print("minimum temp reached")
353             self.current_temperature_accepted_objective_values = [] # reset
354             self.temperature_history.append([self.temperature, self.acceptances_total_count,
355                                               ↪ self.objective_function_evaluation_count])
356             self.Markov_chain_length_current = 0 # restart counter
357             self.acceptances_count_current_chain = 0 # restart counter
358         done = self.get_halt()
359         if done is True:
360             # add final values to make plotting histories easier
361             self.temperature_history.append(
362                 [self.temperature, self.acceptances_total_count,
363                  ↪ self.objective_function_evaluation_count])
364             self.accepted_objective_history.append([self.objective_current,
365                                                     ↪ self.objective_function_evaluation_count])
366             self.acceptances_locations.append(self.objective_function_evaluation_count)
367             self.accepted_x_history.append(self.x_current)
368         return done
369
370     def get_halt(self):
371         """
372         1. first check convergence, converge if over the last evaluation window (5% of total
373         ↪ max function evals), the
374             difference between the maximum and minimum accepted values (within the window) is below
375         ↪ the threshold defined
376             by self.convergence_min_improvement (typically set to 1e-8)
377         2. If the maximum number of function evaluations has been reached (typically set to 10
378         ↪ 000) then end program
379         """
380         if self.objective_function_evaluation_count % self.convergence_evaluation_window == 0:
381             # only make this check every self.convergence_evaluation_window number of
382             ↪ iterations
383             acceptance_locations_array = np.array(self.acceptances_locations)
384             acceptance_indx_over_window = \
385                 np.arange(len(acceptance_locations_array))[
386                     acceptance_locations_array >
387                     self.objective_function_evaluation_count -
388                     ↪ self.convergence_evaluation_window]
389             if len(acceptance_indx_over_window) < 2: # 1 or 0 acceptances within last window
390                 ↪ implies convergence
391                 print("converged")
392                 done = True

```

```

382         return done
383     else: # caclulate difference between max and min over window
384         earliest_acceptance_indx_over_window = acceptance_indx_over_window[0]
385         best_accepted_value_over_recent_window = \
386             np.max(self.accepted_objective_history_array
387                   [earliest_acceptance_indx_over_window:, 0])
388         worst_accepted_value_over_recent_window = \
389             np.min(self.accepted_objective_history_array
390                   [earliest_acceptance_indx_over_window:, 0])
391
392         if best_accepted_value_over_recent_window -
393             ↪ worst_accepted_value_over_recent_window < \
394                 self.convergence_min_improvement:
395             print("converged")
396             done = True
397             return done
398
399 # check if max iter has been reached
400 if self.objective_function_evaluation_count >=
401     ↪ self.objective_function_evaluation_max_count:
402     done = True
403 else:
404     done = False
405 return done
406
407 def update_archive(self, x_new, objective_new):
408     function_archive = [f_archive for x_archive, f_archive in self.archive]
409     dissimilarity = [np.sqrt((x_archive - x_new).T @ (x_archive - x_new)) for x_archive,
410                     ↪ f_archive in self.archive]
411     if min(dissimilarity) > self.archive_minimum_acceptable_dissimilarity: # dissimilar to
412         ↪ all points
413         if len(self.archive) < self.archive_maximum_length: # archive not full
414             self.archive.append((x_new, objective_new))
415         else: # if archive is full
416             if objective_new < min(function_archive):
417                 self.archive[int(np.argmax(function_archive))] = (x_new, objective_new) #
418                 ↪ replace worst solution
419             else: # new solution is close to another
420                 if objective_new < min(function_archive): # objective is lowest yet
421                     most_similar_indx = int(np.argmin(dissimilarity))
422                     self.archive[most_similar_indx] = (x_new, objective_new) # replace most
423                     ↪ similar value
424             else:
425                 similar_and_better = np.array([dissimilarity[i] <
426                 ↪ self.archive_similar_dissimilarity and \
427                     function_archive[i] > objective_new
428                     for i in range(len(self.archive))])
429                 if True in similar_and_better:
430                     self.archive[np.where(similar_and_better == True)[0][0]] = (x_new,
431                     ↪ objective_new)

```

```

423     if self.objective_function_evaluation_count %
424         ↪ (int(self.objective_function_evaluation_max_count/10)) == 0:
425         # sometimes one value can like between 2 others, causing similarity even with the
426         ↪ above loop
427         # clean_archive fixes this
428         # only need to do very rarely
429         self.clean_archive()
430
431     def clean_archive(self):
432         for x_new, y in self.archive:
433             dissimilarity = [np.sqrt((x_archive - x_new).T @ (x_archive - x_new)) for
434                 ↪ x_archive, f_archive in
435                     self.archive]
436             indxs_to_remove = np.where((np.array(dissimilarity) <
437                 ↪ self.archive_minimum_acceptable_dissimilarity) &
438                     (self.archive_f > y)) # remove values that are close,
439                 ↪ with lower objectives
440             indxs_to_remove = indxs_to_remove[0]
441             if len(indxs_to_remove) > 0:
442                 for i, indx_to_remove in enumerate(indxs_to_remove):
443                     # deletions changes indexes so we have to adjust by i each time
444                     del(self.archive[indx_to_remove - i])
445
446     # often it was convenient to store values in lists
447     # however after the optimisation it is more convenient to have
448     # them as arrays, the below property methods are therefore given
449     @property
450     def temperature_history_array(self):
451         return np.array(self.temperature_history)
452
453     @property
454     def archive_x(self):
455         return np.interp(np.array([x_archive for x_archive, f_archive in self.archive]), [-1,
456             ↪ 1], self.x_bounds)
457
458     @property
459     def archive_f(self):
460         return np.array([f_archive for x_archive, f_archive in self.archive])
461
462     @property
463     def accepted_objective_history_array(self):
464         return np.array(self.accepted_objective_history)
465
466     @property
467     def accepted_x_history_array(self):
468         return np.interp(np.array(self.accepted_x_history), [-1, 1], self.x_bounds)
469
470     @property

```

```

466     def objective_history_array(self):
467         return np.array(self.objective_history)
468
469     @property
470     def step_size_matrix_history_array(self):
471         return np.array(self.step_size_matrix_history)
472
473     @property
474     def step_size_update_locations_array(self):
475         return np.array(self.step_size_update_locations)
476
477     @property
478     def probability_of_acceptance_history_array(self):
479         return np.array(self.probability_of_acceptance_history)
480
481     @property
482     def x_history_array(self):
483         return np.interp(np.array(self.x_history), [-1, 1], self.x_bounds)
484
485     @property
486     def alpha_history_array(self):
487         return np.array(self.alpha_history)
488
489     @property
490     def eigenvalue_eigenvector_history(self):
491         theta_history = []
492         eigen_values_history = []
493         for i in range(self.step_size_matrix_history_array.shape[0]):
494             step_size_matrix = self.step_size_matrix_history_array[i, :, :]
495             eigenvalues, eigenvectors = np.linalg.eig(step_size_matrix)
496             thetas = np.arccos(np.eye(self.x_length) @ eigenvectors)
497             min_thetas = np.min(thetas, axis=0)
498             order = np.argsort(-eigenvalues)
499             eigen_values_history.append(list(eigenvalues[order]))
500             theta_history.append(list(min_thetas[order]))
501         return np.array(eigen_values_history), np.array(theta_history)
502
503 if __name__ == "__main__":
504     # simple example run with rana 5D function
505     np.random.seed(0)
506
507     from rana import rana_func
508
509     configuration = {"perturbation_method": "simple",
510                     "x_length": 5,
511                     "x_bounds": (-500, 500),
512                     "annealing_schedule": "simple_exponential_cooling",
513                     "objective_function": rana_func,
514                     "maximum_archive_length": 100,

```

```

515         "archive_minimum_acceptable_dissimilarity": 0.2,
516         "maximum_markov_chain_length": 50,
517         "maximum_function_evaluations": 10000,
518         "step_size_initialisation_fraction_of_range": 0.1,
519         "bound_enforcing_method": "not_clipping",
520         "cholesky_path_length": 5,
521     }
522     np.random.seed(3)
523     rana_2d_chol = SimulatedAnnealing(**configuration)
524     x_result_chol, objective_result_chol = rana_2d_chol.run()
525     print(f"x_result = {x_result_chol} \n objective_result = {objective_result_chol} \n ")
526     print(f"number of function evaluations =
527           ↪ {rana_2d_chol.objective_function_evaluation_count}")
528     print(f"best objective result {rana_2d_chol.objective_history_array.min()}")

```

6.2 Evolution Strategies code

```

1  import numpy as np
2
3
4  class EvolutionStrategy:
5      """
6      x is transformed from the original bounds, to be bounded by (-1, 1)
7      Before outputs are given x is re-transformed back into the original form
8      """
9      def __init__(self, x_length, x_bounds, objective_function,
10                  ↪ archive_minimum_acceptable_dissimilarity=0.1,
11                      parent_number=10,
12                      selection_method="standard_mew_comma_lambda", mutation_method = "simple",
13                      recombination_method="global",
14                      termination_min_abs_difference=1e-6,
15                      maximum_archive_length=None, objective_count_maximum=10000,
16                      mutation_covariance_initialisation_fraction_of_range=0.01,
17                      standard_deviation_clipping_fraction_of_range = 0.05,
18                      bound_enforcing_method="not_clipping",
19                      child_to_parent_ratio=7):
20
21         self.x_length = x_length
22         self.x_bounds = x_bounds
23         self.bound_enforcing_method = bound_enforcing_method
24         self.x_range = 2 # after being transformed to be between -1 and 1
25         self.objective_function_raw = objective_function
26         self.selection_method = selection_method
27         self.mutation_method = mutation_method
28         self.recombination_method = recombination_method
29         self.termination_min_abs_difference = termination_min_abs_difference
30
31         # prevent standard deviations from becoming too large (clip relative to size of range)

```

```

30 self.standard_deviation_clipping_fraction_of_range =
    ↪ standard_deviation_clipping_fraction_of_range
31 self.parent_number = parent_number
32 self.offspring_number = parent_number * child_to_parent_ratio
33
34 if mutation_method == "complex":
35     # Mutation parameters - taken from slides, recommended by (Schwefel 1987)
36     self.mutation_tau = 1/np.sqrt(2*np.sqrt(self.x_length))
37     self.mutation_tau_dash = 1/np.sqrt(2*self.x_length)
38     self.mutation_Beta = 0.0873
39     self.offspring_mutation_standard_deviations = \
40         np.ones((self.offspring_number, self.x_length)) * \
41         mutation_covariance_initialisation_fraction_of_range * self.x_range
42     self.parent_mutation_standard_deviations = \
43         np.ones((self.parent_number, self.x_length)) * \
44         mutation_covariance_initialisation_fraction_of_range * self.x_range
45
46     self.offspring_rotation_matrices = np.broadcast_to(np.eye(self.x_length),
    ↪ (self.offspring_number, self.x_length, self.x_length))
47
48     self.make_covariance_matrix()
49     # just slice children for initialisation
50     self.parent_rotation_matrices =
    ↪ self.offspring_rotation_matrices[0:self.parent_number, :, :]
51     self.parent_covariance_matrices =
    ↪ self.offspring_covariance_matrices[0:self.parent_number, :, :]
52
53
54 elif mutation_method == "simple":
55     self.standard_deviation_simple =
    ↪ mutation_covariance_initialisation_fraction_of_range*self.x_range
56
57 elif mutation_method == "diagonal":
58     self.mutation_tau = 1 / np.sqrt(2 * np.sqrt(self.x_length))
59     self.mutation_tau_dash = 1 / np.sqrt(2 * self.x_length)
60     self.offspring_mutation_standard_deviations = \
61         np.ones((self.offspring_number, self.x_length)) * \
62         mutation_covariance_initialisation_fraction_of_range * self.x_range
63     self.parent_mutation_standard_deviations = np.ones((self.parent_number,
    ↪ self.x_length)) * mutation_covariance_initialisation_fraction_of_range *
    ↪ self.x_range
64
65 # initialise parents and offspring
66 # zeros aren't ever used, just specifies the shapes of the arrays
67 self.parents = np.zeros((self.parent_number, self.x_length))
68 self.parent_objectives = np.zeros(self.parent_number)
69 self.offspring = np.zeros((self.offspring_number, self.x_length))
70 self.offspring_objectives = np.zeros(self.offspring_number)
71

```

```

72
73     # initialise archive and parameters determining how archive is managed
74     self.archive = []    # list of (x, objective value) tuples
75     self.archive_maximum_length = maximum_archive_length    # If none then don't store, as
76     ↪ slows program down slightly
77     self.archive_minimum_acceptable_dissimilarity =
78     ↪ archive_minimum_acceptable_dissimilarity
79     self.archive_similar_dissimilarity = archive_minimum_acceptable_dissimilarity
80
81     # initialise histories and counters
82     # these are useful for inspecting the performance of the algorithm after a run
83     # and are used within the program (e.g. markov chain length)
84     self.parent_objective_history = []
85     self.parent_x_history = []
86     self.parent_standard_deviation_history = []
87     self.offspring_objective_history = []
88     self.offspring_x_history = []
89     self.parent_covariance_determinant_history = []
90     self.offspring_covariance_determinant_history = []
91     self.objective_function_evaluation_count = 0 # initialise
92     self.generation_number = 0
93     self.objective_function_evaluation_max_count = objective_count_maximum
94
95 def objective_function(self, x):
96     """
97     Wrapper for the objective function calls, adding some extra functionality
98     """
99     # increment by one everytime objective function is called
100     self.objective_function_evaluation_count += 1
101     # interpolation done here to pass the objective function x correctly interpolated
102     x_interp = np.interp(x, [-1, 1], self.x_bounds)
103     result = self.objective_function_raw(x_interp)
104     return result
105
106 def run(self):
107     """
108     This function run's the major steps of the Evolution Strategy algorithm
109     # major steps in the algorithm are surrounded with a #*****
110     # other parts of this function are more organisation (e.g. storing histories)
111     """
112     self.initialise_random_population()
113     while True: # loop until termination criteria is reached
114         self.generation_number += 1
115         # ***** Selection *****
116         self.select_parents()
117         # *****
118         if self.archive_maximum_length is not None: # if archive is None, then don't store
119             for x, objective in zip(self.parents, self.parent_objectives): # update
120                 ↪ archive

```



```

118         self.update_archive(x, objective)
119     self.parent_objective_history.append(self.parent_objectives)
120     self.parent_x_history.append(self.parents)
121     if self.mutation_method == "diagonal":
122         self.parent_standard_deviation_history.append\
123             (self.parent_mutation_standard_deviations)
124         self.parent_covariance_determinant_history.append(
125             np.prod(self.parent_mutation_standard_deviations, axis=1))
126         self.offspring_covariance_determinant_history.append(
127             np.prod(self.offspring_mutation_standard_deviations, axis=1))
128     elif self.mutation_method == "complex":
129         self.parent_standard_deviation_history.append(
130             self.parent_mutation_standard_deviations)
131         self.parent_covariance_determinant_history.append(
132             np.linalg.det(self.parent_covariance_matrices))
133         self.offspring_covariance_determinant_history.append(
134             np.linalg.det(self.offspring_covariance_matrices))
135
136     # ***** Check for convergence/termination *****
137     # ensure termination before 10000 iterations
138     if self.objective_function_evaluation_count >
139         ↪ self.objective_function_evaluation_max_count-self.offspring_number:
140         print("max total iterations")
141         break
142
143     # termination criteria
144     if max(self.parent_objectives) - min(self.parent_objectives) <
145         ↪ self.termination_min_abs_difference:
146         print("converged")
147         break
148
149     # ***** Create Offspring *****
150     # I.e. perform recombination and mutation to create offspring
151     self.create_new_offspring()
152     # *****
153     self.offspring_objective_history.append(self.offspring_objectives)
154     self.offspring_x_history.append(self.offspring)
155
156     best_x = self.parents[np.argmin(self.parent_objectives), :]
157     best_objective = min(self.parent_objectives)
158     return np.interp(best_x, [-1, 1], self.x_bounds), best_objective
159
160 def initialise_random_population(self):
161     self.offspring = np.random.uniform(low=-1, high=1, size=(self.offspring_number,
162         ↪ self.x_length))
163     self.offspring_objectives =
164     ↪ np.squeeze(np.apply_along_axis(func1d=self.objective_function, arr=self.offspring,
165     ↪ axis=1))
166     if self.selection_method == "elitist": # require pool including parents for
167     ↪ select_parents function in this case

```

```

162         self.parents = np.random.uniform(low=-1, high=1,
163                                           size=(self.parent_number, self.x_length))
164         self.parent_objectives = np.apply_along_axis(funcid=self.objective_function,
165                                                       ↪ arr=self.parents, axis=1)
166
167     def select_parents(self):
168         if self.selection_method == "standard_mew_comma_lambda":
169             # choose top values in linear time
170             # np.argmax doesn't sort top values amongst themselves so is computationally
171             ↪ faster
172             pool_objectives = self.offspring_objectives
173             pool = self.offspring
174             if self.mutation_method == "diagonal":
175                 pool_standard_deviations = self.offspring_mutation_standard_deviations
176             if self.mutation_method == "complex":
177                 pool_standard_deviations = self.offspring_mutation_standard_deviations
178                 pool_rotation_matrices = self.offspring_rotation_matrices
179                 pool_covariance_matrices = self.offspring_covariance_matrices
180         else:
181             assert self.selection_method == "elitist"
182             # create pool selected from
183             pool_objectives = np.zeros(self.parent_number + self.offspring_number)
184             pool_objectives[0:self.offspring_number] = self.offspring_objectives
185             pool_objectives[self.offspring_number:] = self.parent_objectives
186             pool = np.zeros((self.offspring_number + self.parent_number, self.x_length))
187             pool[0:self.offspring_number, :] = self.offspring
188             pool[self.offspring_number:, :] = self.parents
189             if self.mutation_method == "diagonal":
190                 pool_standard_deviations = np.zeros((self.offspring_number +
191                                                       ↪ self.parent_number, self.x_length))
192                 pool_standard_deviations[0:self.offspring_number, :] =
193                 ↪ self.offspring_mutation_standard_deviations
194                 pool_standard_deviations[self.offspring_number:, :] =
195                 ↪ self.parent_mutation_standard_deviations
196             if self.mutation_method == "complex":
197                 pool_standard_deviations = np.zeros((self.offspring_number +
198                                                       ↪ self.parent_number, self.x_length))
199                 pool_standard_deviations[0:self.offspring_number, :] =
200                 ↪ self.offspring_mutation_standard_deviations
201                 pool_standard_deviations[self.offspring_number:, :] =
202                 ↪ self.parent_mutation_standard_deviations
203                 pool_rotation_matrices = np.zeros((self.offspring_number + self.parent_number,
204                                                       ↪ self.x_length, self.x_length))
205                 pool_rotation_matrices[0:self.offspring_number, :, :] =
206                 ↪ self.offspring_rotation_matrices
207                 pool_rotation_matrices[self.offspring_number:, :, :] =
208                 ↪ self.parent_rotation_matrices
209                 pool_covariance_matrices = np.zeros((self.offspring_number +
210                                                       ↪ self.parent_number, self.x_length, self.x_length))

```

```

199         pool_covariance_matrices[0:self.offspring_number, :, :] =
200             ↪ self.offspring_covariance_matrices
201
202         pool_covariance_matrices[self.offspring_number:, :, :] =
203             ↪ self.parent_covariance_matrices
204
205     new_parent_indxs = np.argpartition(pool_objectives,
206         ↪ self.parent_number)[:self.parent_number]
207     self.parents = pool[new_parent_indxs, :]
208     self.parent_objectives = pool_objectives[new_parent_indxs]
209
210     if self.mutation_method == "diagonal":
211         self.parent_mutation_standard_deviations =
212             ↪ pool_standard_deviations[new_parent_indxs, :]
213     elif self.mutation_method == "complex":
214         self.parent_mutation_standard_deviations =
215             ↪ pool_standard_deviations[new_parent_indxs, :]
216         self.parent_rotation_matrices = pool_rotation_matrices[new_parent_indxs, :, :]
217         self.parent_covariance_matrices = pool_covariance_matrices[new_parent_indxs, :, :]
218
219
220 def create_new_offspring(self):
221     """
222     Recombination and Mutation
223     """
224     ***** Recombination *****
225     # global discrete recombination for control parameters
226     # global intermediate recombination for strategy parameters
227     if self.recombination_method == "global":
228         # for each element in each child, inherit from a random parent
229         child_recombination_indxs = np.random.choice(self.parent_number, replace=True,
230             size = (self.offspring_number, self.x_length))
231         offspring_pre_mutation = self.parents[child_recombination_indxs,
232             ↪ np.arange(self.x_length)]
233         if self.mutation_method == "diagonal" or self.mutation_method == "complex":
234             child_strategy_recombination_indxs = np.random.choice(self.parent_number,
235                 ↪ replace=True,
236
237                 size=(self.offspring_number, self.x_length, 2))
238             offspring_pre_mutation_standard_deviation = \
239                 0.5*self.parent_mutation_standard_deviations[
240                     child_strategy_recombination_indxs[:, :, 0], np.arange(self.x_length)]
241                 ↪ +\
242                 0.5*self.parent_mutation_standard_deviations[
243                     child_strategy_recombination_indxs[:, :, 1], np.arange(self.x_length)]
244         if self.mutation_method == "complex":
245             child_strategy_recombination_indxs = \
246                 np.random.choice(self.parent_number, replace=True,
247                     ↪ size=(self.offspring_number*self.x_length*self.x_length, 2))
248             slices1 = np.broadcast_to(np.arange(self.x_length)[np.newaxis, :],
249                 ↪ (self.x_length, self.x_length)).flatten()

```

```

239         slices1 = np.broadcast_to(slices1[np.newaxis, :], (self.offspring_number,
240             ↪ self.x_length**2)).flatten()
241         slices2 = np.broadcast_to(np.arange(self.x_length)[: , np.newaxis],
242             ↪ (self.x_length, self.x_length)).flatten()
243         slices2 = np.broadcast_to(slices2[np.newaxis, :], (self.offspring_number,
244             ↪ self.x_length**2)).flatten()
245         offspring_pre_mutation_rotation_matrix = \
246             np.reshape(0.5 * self.parent_rotation_matrices[
247                 child_stratergy_recombination_indxs[:, 0], slices1,
248                 slices2 ] + \
249                 0.5 * self.parent_rotation_matrices[
250                     child_stratergy_recombination_indxs[:, 1], slices1,
251                     slices2], (self.offspring_number, self.x_length, self.x_length))
252
253     ##### Mutation #####
254     if self.mutation_method == "simple":
255         u_random_sample = np.random.normal(loc=0, scale=self.standard_deviation_simple,
256             size=offspring_pre_mutation.shape)
257         x_new = offspring_pre_mutation + u_random_sample
258         if self.bound_enforcing_method == "clipping":
259             x_new = np.clip(x_new, -1, 1)
260         else:
261             while np.max(x_new) > 1 or np.min(x_new) < -1:
262                 indxs_breaking_bounds = np.where((x_new > 1) + (x_new < -1) == 1)
263                 u_random_sample = np.random.normal(loc=0,
264                     ↪ scale=self.standard_deviation_simple,
265                     size=indxs_breaking_bounds[0].size)
266                 x_new[indxs_breaking_bounds ] =
267                     ↪ offspring_pre_mutation[indxs_breaking_bounds ] + u_random_sample
268
269     elif self.mutation_method == "diagonal": # non spherical covariance
270         self.offspring_mutation_standard_deviations = \
271             offspring_pre_mutation_standard_deviation * \
272             np.exp(self.mutation_tau_dash*np.broadcast_to(
273                 np.random.normal(0, 1, size=(self.offspring_number, 1)),
274                 ↪ self.offspring_mutation_standard_deviations.shape)
275                 + self.mutation_tau*np.random.normal(0, 1,
276                 ↪ size=self.offspring_mutation_standard_deviations.shape))
277         self.offspring_mutation_standard_deviations = \
278             np.clip(self.offspring_mutation_standard_deviations,
279                 1e-8, self.standard_deviation_clipping_fraction_of_range*self.x_range)
280
281         u_random_sample = np.random.normal(loc=0,
282             ↪ scale=self.offspring_mutation_standard_deviations,
283             size=offspring_pre_mutation.shape)
284         x_new = offspring_pre_mutation + u_random_sample
285         if self.bound_enforcing_method == "clipping":
286             x_new = np.clip(x_new, -1, 1)

```

```

280         else:
281             while np.max(x_new) > 1 or np.min(x_new) < -1:
282                 indxs_breaking_bounds = np.where((x_new > 1) + (x_new < -1) == 1)
283                 u_random_sample = \
284                     np.random.normal(loc=0,
285                                     ↪ scale=self.offspring_mutation_standard_deviations
286                                     [indxs_breaking_bounds],size=indxs_breaking_bounds[0].size)
287                 x_new[indxs_breaking_bounds] =
288                     ↪ offspring_pre_mutation[indxs_breaking_bounds] + u_random_sample
289
290     if self.mutation_method == "complex":
291         self.offspring_mutation_standard_deviations = \
292             offspring_pre_mutation_standard_deviation * \
293             np.exp(self.mutation_tau_dash*np.broadcast_to(
294                 np.random.normal(0, 1, size=(self.offspring_number, 1)),
295                 ↪ self.offspring_mutation_standard_deviations.shape)
296             + self.mutation_tau*np.random.normal(0, 1,
297             ↪ size=self.offspring_mutation_standard_deviations.shape))
298
299         self.offspring_mutation_standard_deviations = \
300             np.clip(self.offspring_mutation_standard_deviations, 1e-8,
301                     self.standard_deviation_clipping_fraction_of_range * self.x_range)
302
303         self.offspring_rotation_matrices = offspring_pre_mutation_rotation_matrix +
304             ↪ self.mutation_Beta * np.random.normal(0, 1,
305             ↪ size=(self.offspring_rotation_matrices.shape))
306
307         for i in range(self.offspring_number):
308             self.offspring_rotation_matrices[i, :, :] =
309                 ↪ np.tril(self.offspring_rotation_matrices[i, :, :], k=-1) - np.tril(
310                     self.offspring_rotation_matrices[i, :, :], k=-1).T      # make symmetric
311         self.make_covariance_matrix()
312         for i in range(self.offspring_number):
313             covariance_matrix =
314                 ↪ self.make_positive_definite(self.offspring_covariance_matrices[i, :, :])
315             self.offspring[i, :] = offspring_pre_mutation[i, :] +
316                 ↪ np.random.multivariate_normal(mean=np.zeros(self.x_length),
317                 ↪ cov=covariance_matrix)
318
319         if self.bound_enforcing_method == "clipping":
320             self.offspring[i, :] = np.clip(self.offspring[i, :] , -1, 1)
321         else:
322             while np.max(self.offspring[i, :]) > 1 or np.min(self.offspring[i, :]) <
323                 ↪ -1:
324                 self.offspring[i, :] = offspring_pre_mutation[i, :] +
325                     ↪ np.random.multivariate_normal(mean=np.zeros(self.x_length),
326                     ↪ cov=covariance_matrix)
327
328     if self.mutation_method != "complex":
329         self.offspring = x_new
330     self.offspring_objectives = np.squeeze(
331         np.apply_along_axis(func1d=self.objective_function, arr=self.offspring, axis=1))

```

```

316
317 def make_positive_definite(self, matrix, i=1):
318     try:
319         np.linalg.cholesky(matrix)
320         return matrix
321     except:
322         if i > 10:
323             raise Exception("matrix unable to be made positive definite")
324         matrix += np.eye(self.x_length) * 1e-6 * 10**i
325         return self.make_positive_definite(matrix, i=i+1)
326
327
328 def update_archive(self, x_new, objective_new):
329     if len(self.archive) == 0: # if empty then initialise with the first value
330         self.archive.append((x_new, objective_new))
331     function_archive = [f_archive for x_archive, f_archive in self.archive]
332     dissimilarity = [np.sqrt((x_archive - x_new).T @ (x_archive - x_new)) for x_archive,
333                     ↪ f_archive in self.archive]
334     if min(dissimilarity) > self.archive_minimum_acceptable_dissimilarity: # dissimilar to
335     ↪ all points
336         if len(self.archive) < self.archive_maximum_length: # archive not full
337             self.archive.append((x_new, objective_new))
338         else: # if archive is full
339             if objective_new < min(function_archive):
340                 self.archive[int(np.argmax(function_archive))] = (x_new, objective_new) #
341                 ↪ replace worst solution
342             else: # new solution is close to another
343                 if objective_new < min(function_archive): # objective is lowest yet
344                     most_similar_idx = int(np.argmin(dissimilarity))
345                     self.archive[most_similar_idx] = (x_new, objective_new) # replace most
346                     ↪ similar value
347                 else:
348                     similar_and_better = np.array([dissimilarity[i] <
349                     ↪ self.archive_similar_dissimilarity and \
350                                     function_archive[i] > objective_new
351                                     for i in range(len(self.archive))])
352                     if True in similar_and_better:
353                         self.archive[np.where(similar_and_better == True)[0][0]] = (x_new,
354                         ↪ objective_new)
355             if self.generation_number % 10 == 0:
356                 # sometimes one value can like between 2 others, causing similarity even with the
357                 ↪ above loop
358                 # clean_archive fixes this
359                 # only need to do very rarely
360                 # slows down program a lot, so only perform when we need to visualise 2D problem
361                 self.clean_archive()
362
363 def clean_archive(self):
364     # first remove repeats

```

```

358     for x_new, y in self.archive:
359         dissimilarity = [np.sqrt((x_archive - x_new).T @ (x_archive - x_new)) for
360             ↪ x_archive, f_archive in
361                 self.archive]
362         indxs_to_remove = np.where(np.array(dissimilarity) == 0) # remove values that are
363             ↪ close, with lower objectives
364         indxs_to_remove = indxs_to_remove[0]
365         if len(indxs_to_remove) > 0:
366             indxs_to_remove = indxs_to_remove[1:] # remove all but the first copy
367             for i, indx_to_remove in enumerate(indxs_to_remove):
368                 # deletions changes indexes so we have to adjust by i each time
369                 del (self.archive[indx_to_remove - i])
370
371     # then remove overly similar
372     for x_new, y in self.archive:
373         dissimilarity = [np.sqrt((x_archive - x_new).T @ (x_archive - x_new)) for
374             ↪ x_archive, f_archive in
375                 self.archive]
376         indxs_to_remove = np.where((np.array(dissimilarity) <
377             ↪ self.archive_minimum_acceptable_dissimilarity) &
378             (self.archive_f > y)) # remove values that are close,
379             ↪ with lower objectives
380         indxs_to_remove = indxs_to_remove[0]
381         if len(indxs_to_remove) > 0:
382             for i, indx_to_remove in enumerate(indxs_to_remove):
383                 # deletions changes indexes so we have to adjust by i each time
384                 del (self.archive[indx_to_remove - i])
385
386     def make_covariance_matrix(self):
387         sigma_i = np.zeros((self.offspring_number, self.x_length, self.x_length))
388         sigma_j = np.zeros((self.offspring_number, self.x_length, self.x_length))
389         for offspring_number in range(self.offspring_number):
390             stds = self.offspring_mutation_standard_deviations[offspring_number, :]
391             sigma_i[offspring_number, np.arange(self.x_length), :] = stds
392             sigma_j[offspring_number, :, np.arange(self.x_length)] = stds
393         self.offspring_covariance_matrices = np.tan(2 * self.offspring_rotation_matrices) *
394             ↪ (sigma_i**2 - sigma_j**2) * 1/2
395         self.offspring_covariance_matrices = np.clip(self.offspring_covariance_matrices,
396             ↪ -np.minimum(sigma_i, sigma_j), np.minimum(sigma_i, sigma_j))
397         self.offspring_covariance_matrices[:, np.arange(self.x_length),
398             ↪ np.arange(self.x_length)] = self.offspring_mutation_standard_deviations
399
400     # often it was conventient to store values in lists
401     # however after the optimisation it is more convenient to have
402     # them as arrays, the below property methods are therefore given
403     @property
404     def parent_objective_history_array(self):
405         return np.array(self.parent_objective_history)

```

```

399 @property
400 def offspring_objective_history_array(self):
401     return np.array(self.offspring_objective_history)
402
403 @property
404 def parent_standard_deviation_history_array(self):
405     return np.array(self.parent_standard_deviation_history)
406
407 @property
408 def parent_covariance_determinant_history_array(self):
409     return np.array(self.parent_covariance_determinant_history)
410
411 @property
412 def offspring_covariance_determinant_history_array(self):
413     return np.array(self.offspring_covariance_determinant_history)
414
415 @property
416 def offspring_x_history_array(self):
417     return np.interp(np.array(self.offspring_x_history), [-1, 1], self.x_bounds)
418
419 @property
420 def parent_x_history_array(self):
421     return np.interp(np.array(self.parent_x_history), [-1, 1], self.x_bounds)
422
423 @property
424 def archive_x(self):
425     return np.interp(np.array([x_archive for x_archive, f_archive in self.archive]), [-1,
426     ↪ 1], self.x_bounds)
427
428 @property
429 def archive_f(self):
430     return np.array([f_archive for x_archive, f_archive in self.archive])
431
432 if __name__ == "__main__":
433     # example run on the 5 D rana problem
434     from rana import rana_func
435     Comp_config = {"objective_function": rana_func,
436                   "x_bounds": (-500, 500),
437                   "x_length": 5,
438                   "parent_number": 10,
439                   "child_to_parent_ratio": 7,
440                   "bound_enforcing_method": "not_clipping",
441                   "selection_method": "standard_mew_comma_lambda",
442                   "standard_deviation_clipping_fraction_of_range": 0.01,
443                   "mutation_covariance_initialisation_fraction_of_range": 0.01,
444                   "mutation_method": "complex",
445                   "termination_min_abs_difference": 1e-6,
446                   "maximum_archive_length": 20}
447
448     random_seed = 1

```



```

447     np.random.seed(random_seed)
448     x_max = 500
449     x_min = -x_max
450     evo_comp = EvolutionStrategy(**Comp_config)
451     x_result, objective_result = evo_comp.run()
452     print(f"x_result = {x_result} \n objective_result = {objective_result}\n\n\n\
453           number of objective_evaluations is {evo_comp.objective_function_evaluation_count}\
454           number of generations is {evo_comp.generation_number}")

```

6.3 Random search code

```

1  import numpy as np
2  from rana import rana_func
3  class random_search:
4      def __init__(self, x_length, x_bounds= (-500, 500), objective_function=rana_func,
5          ↪ n_evaluations=10000):
6          self.x_length = x_length
7          self.x_bounds = x_bounds
8          self.objective_function = objective_function
9          self.n_evaluations = n_evaluations
10
11     def run(self):
12         self.all_points = np.random.uniform(low=self.x_bounds[0], high=self.x_bounds[1],
13             ↪ size=(self.n_evaluations, self.x_length))
14         self.objectives = np.apply_along_axis(func1d=self.objective_function,
15             ↪ arr=self.all_points, axis=1)
16         return self.objectives.min(), self.all_points[np.argmin(self.objectives), :]

```